

Fallacies of Distributed computing

and their effect on Cloud Solutions



Amo Abeyaratne (5/3/19)

What is this session about?

- What to think about when building data pipelines in Cloud platforms
- Managed is a big deal - but can we take it for granted?
- Best practices for architecting Cloud Data platforms (and how some of the veteran cloud customers deal with it)

This is the opposite of a cloud sales pitch: But don't forget the following..



Edy S. Liongosari

@edyliong

Follow



“@joeweinman: @simoncrosby at
#structureconf Cloud outages are like plane
crashes: lots of publicity but still safer than
driving yourself”

12:08 PM - 22 Jun 2011




“just as flying is still the safest form of travel in terms of number of fatalities per miles travelled, cloud computing services provide as high—or higher—levels of availability and reliability as enterprise data centers.”

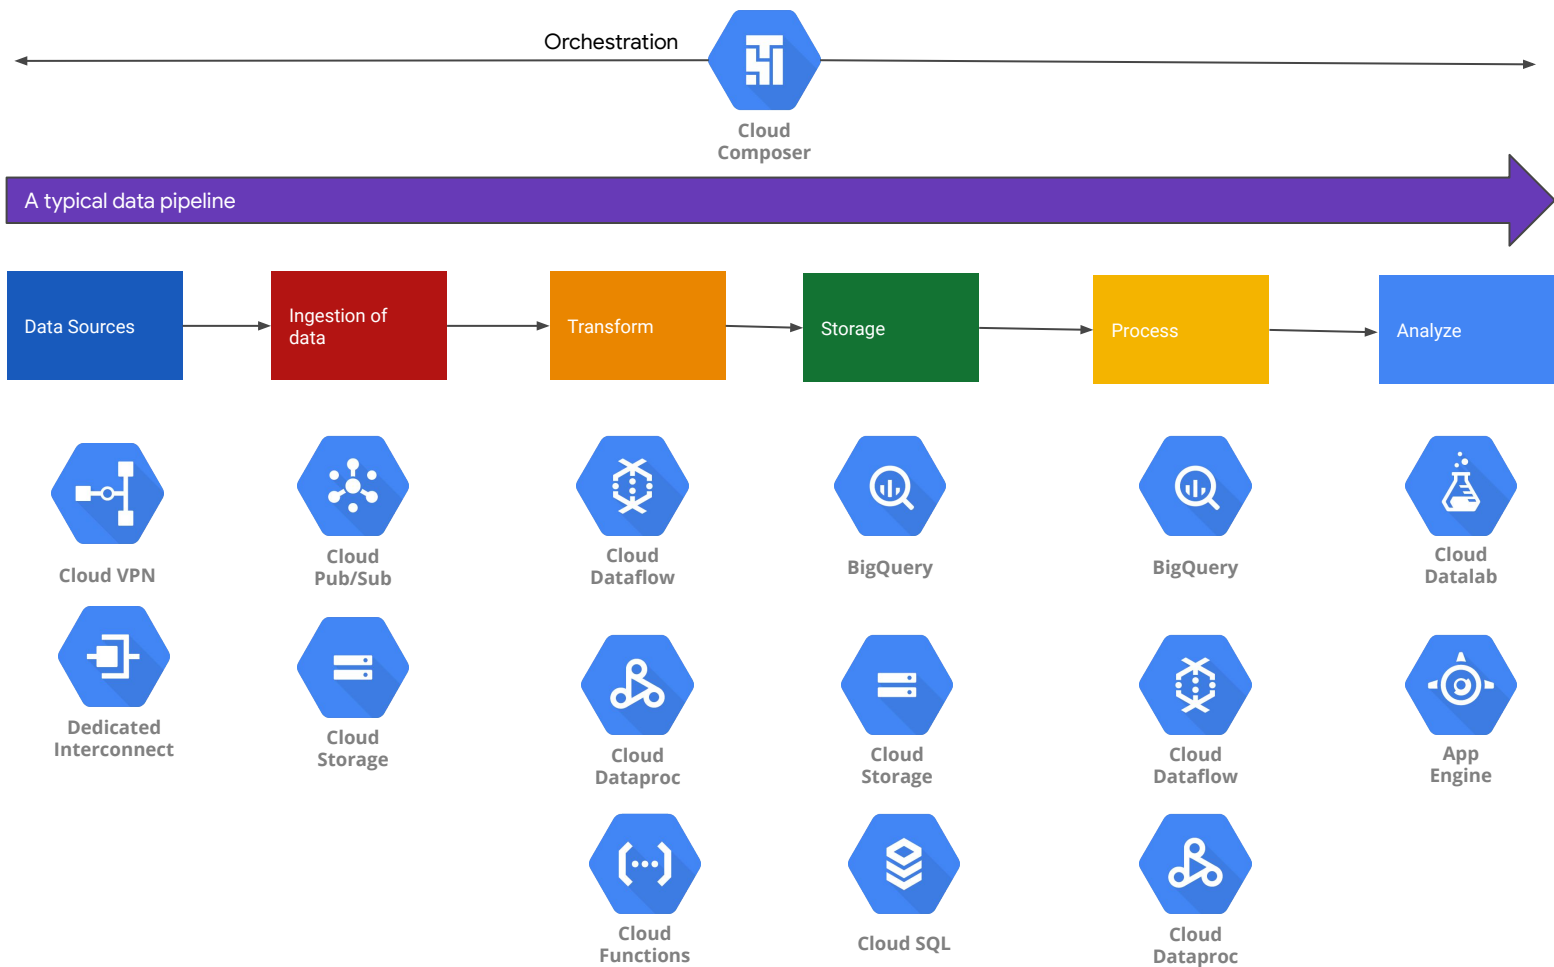
Let's look at a typical data pipeline



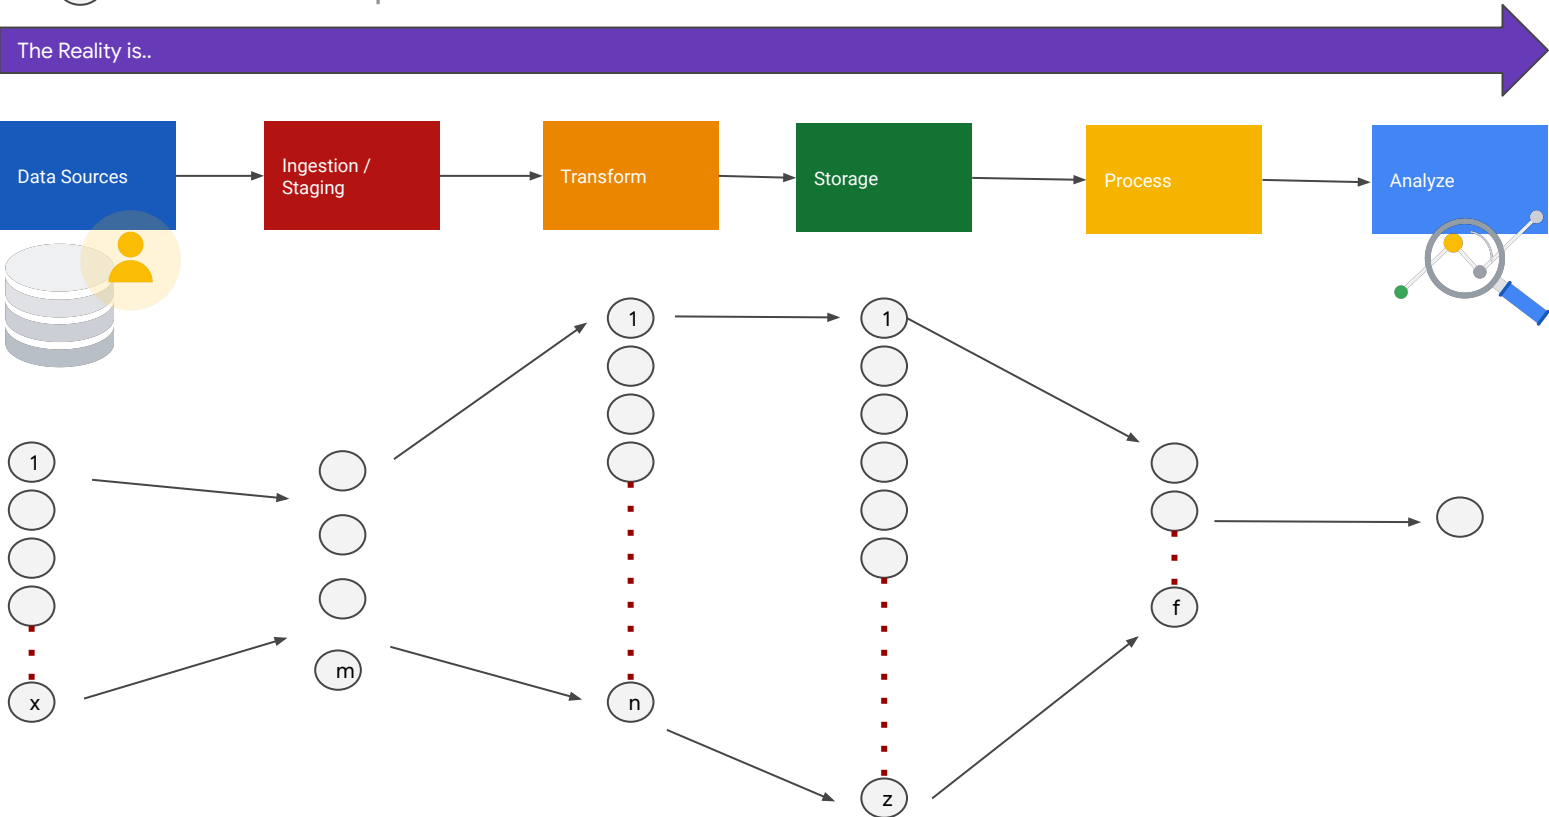
How would you build it in
the cloud?

A typical data pipeline





If ○ = Network component / Node




What is a Distributed System?

What is a Distributed System?

A distributed system is a system whose components are located on different networked computers, which communicate and coordinate their actions by passing messages to one another.


Source: https://en.wikipedia.org/wiki/Distributed_computing

(a)




Proprietary + Confidential

(b)



(c)



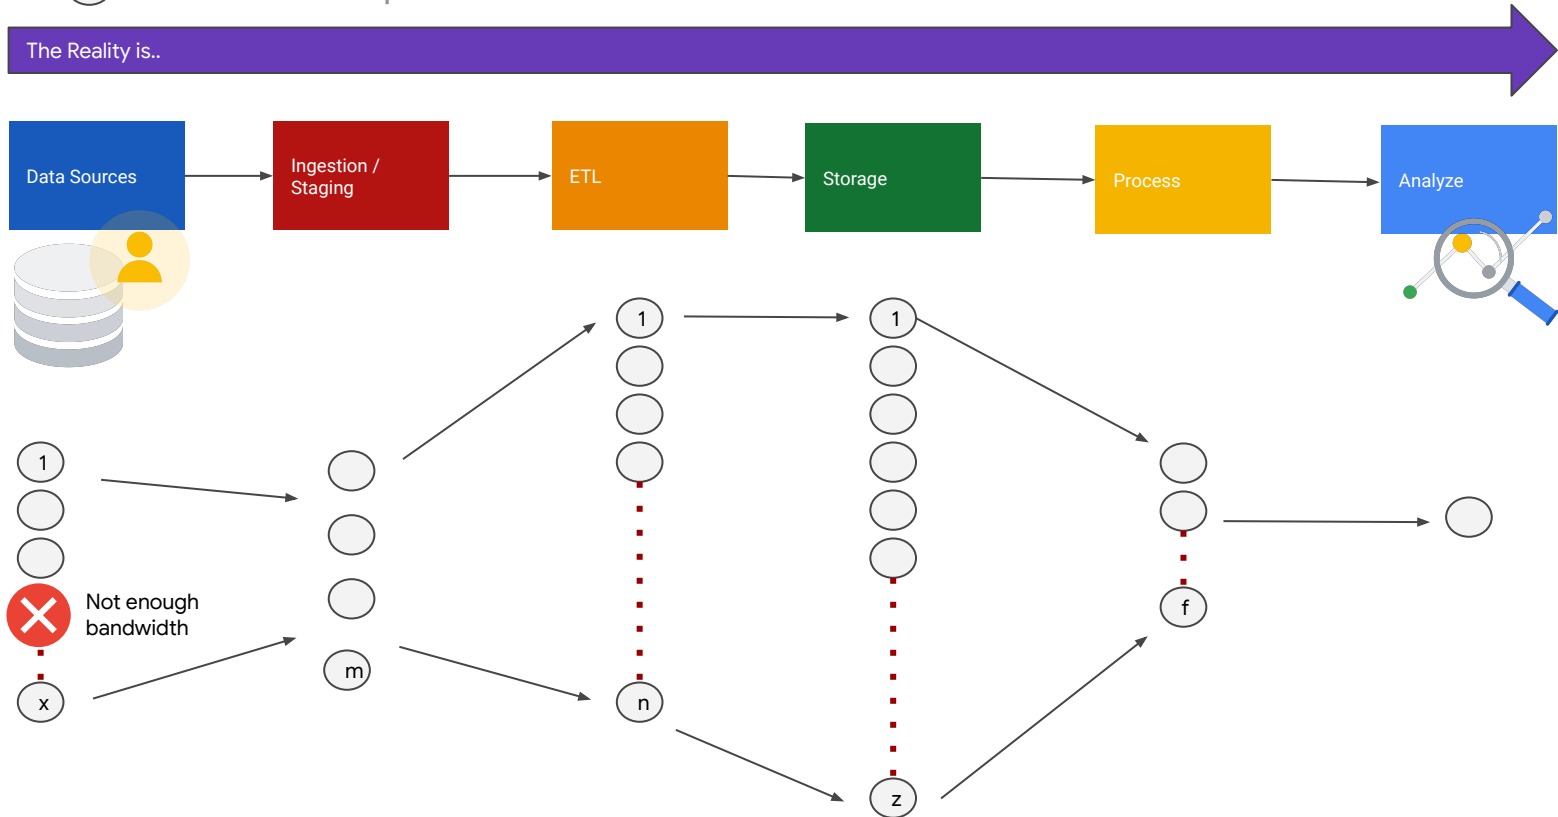
Google

What can fail? }

Anything that can fail, will.

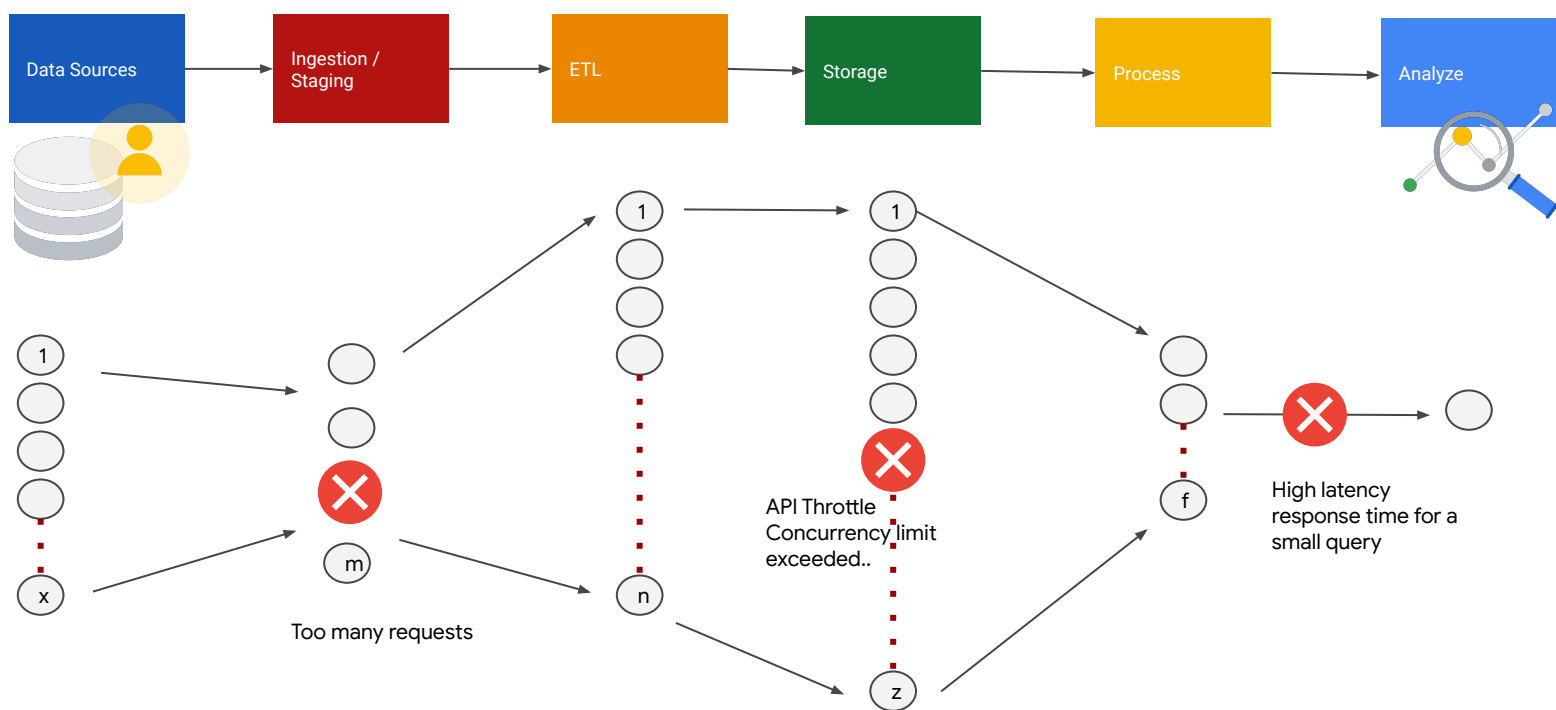
If ○ = Network component / Node

The Reality is..



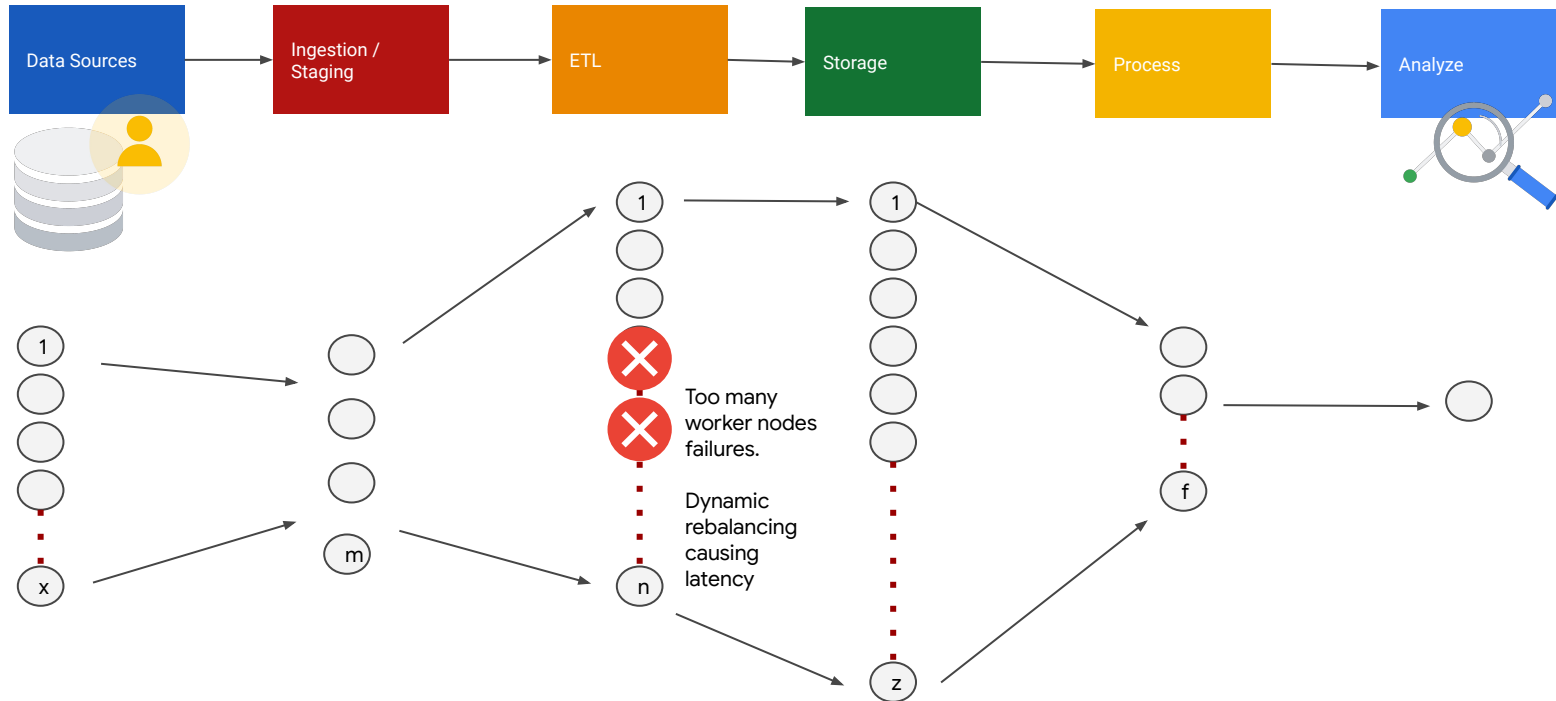
If ○ = Network component / Node

The Reality is..



If ○ = Network component / Node

The Reality is.... Things will always fail, at cloud scale it fails more frequently..



What do we
take for
granted here?



"The fallacies of Distributed
Computing"


The fallacies of Distributed Computing

1. The network is reliable.
2. Latency is zero.
3. Bandwidth is infinite.
4. The network is secure.
5. Topology doesn't change.
6. There is one administrator.
7. Transport cost is zero.
8. The network is homogeneous.



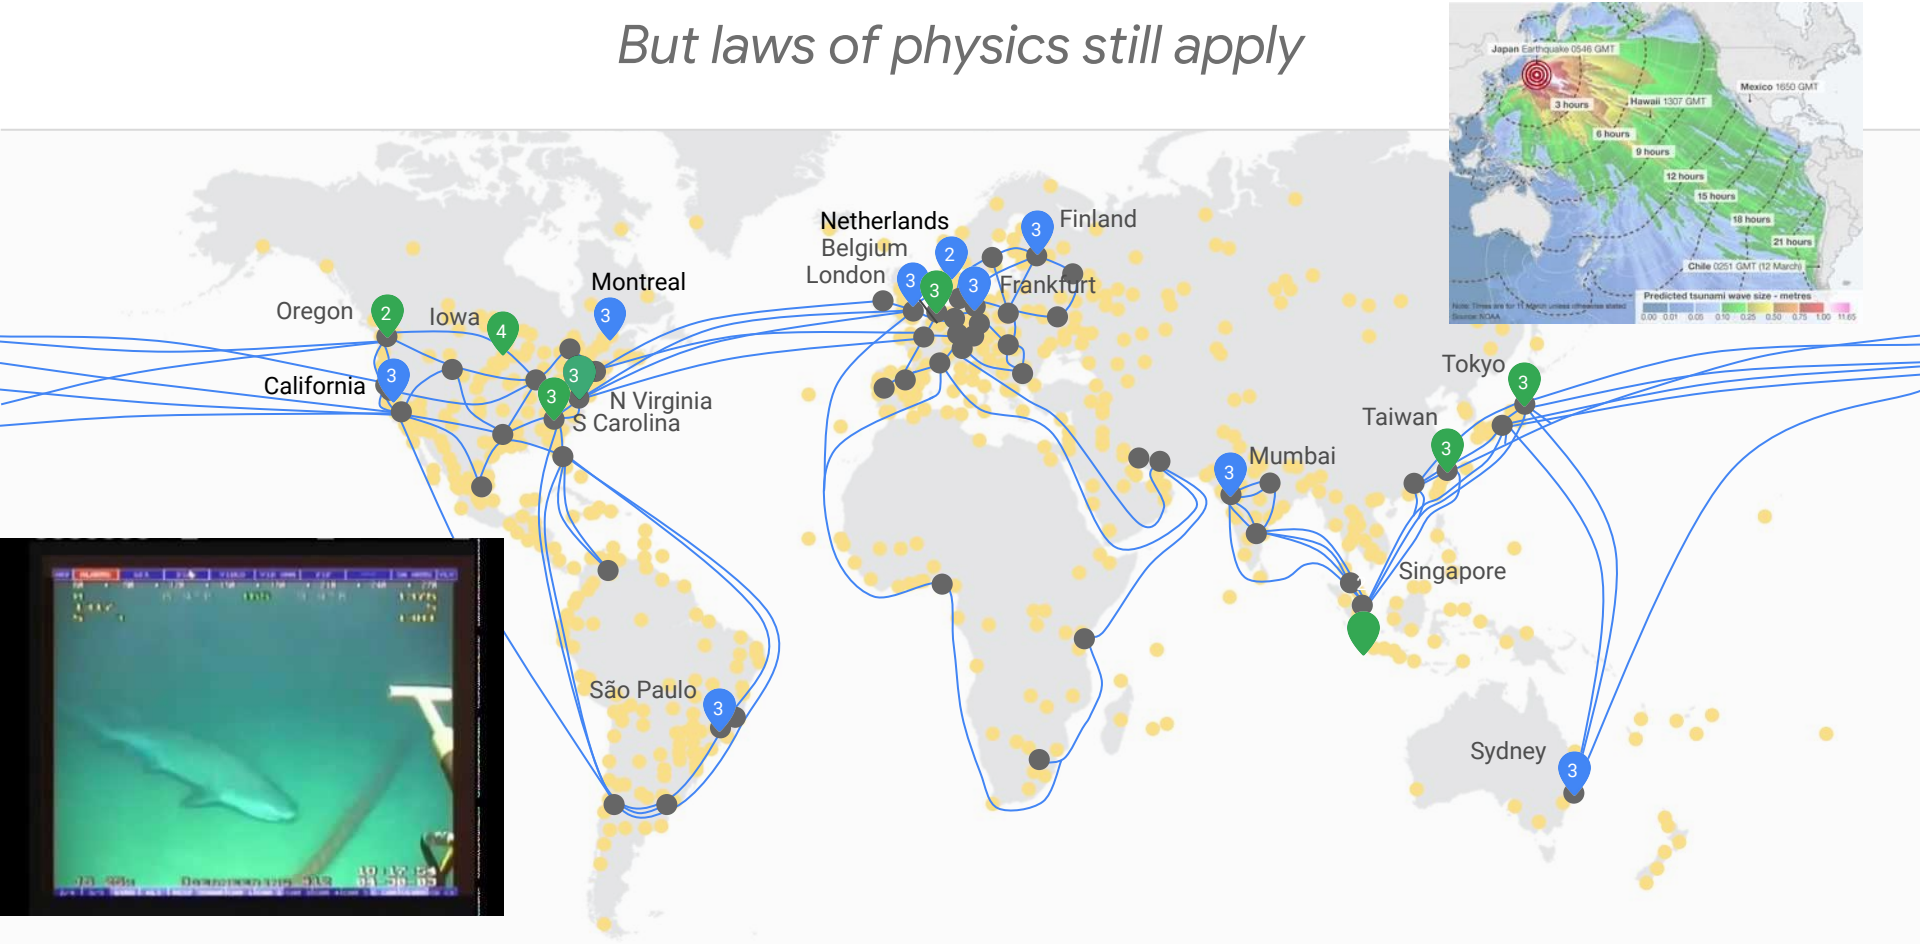
Peter Deutsch - Sun
Microsystems, 1994

How much time you would spend uploading data?



Hundreds of thousands of miles of fiber optic cable
connecting all of our datacenter regions and 100+ points of presence.

But laws of physics still apply



So, what should
we do? }

Design for failure :)

Start your solution design with the end user in mind..

*<https://landing.google.com/sre/books/> (if interested to learn more about Google SRE philosophy)

The fallacies of Distributed Computing

Meets

Design principles for Cloud Solutions

1. The network is reliable.

2. Latency is zero.

3. Bandwidth is infinite.

4. The network is secure.

5. Topology doesn't change.

6. There is one administrator.

7. Transport cost is zero.

8. The network is homogeneous.

1. Design for failure

2. Loose coupling

3. Implement “Elasticity”

4. Build Security in every layer

5. Don't fear constraints

6. Think Parallel

7. Leverage different storage options

This doesn't matter anymore, we have achieved a good level of abstraction

* These best practices for Architecting in Cloud was introduced by AWS in 2011.

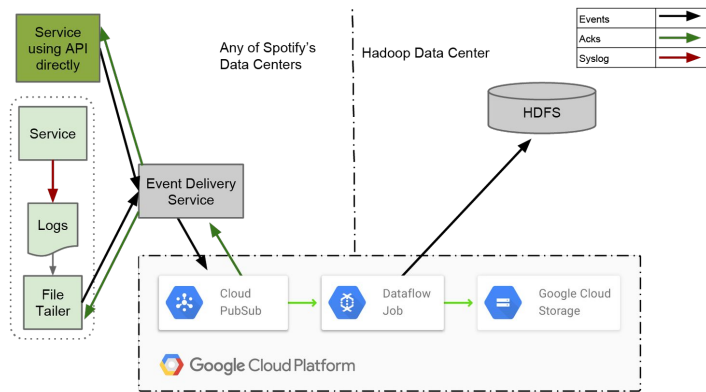
<https://aws.amazon.com/whitepapers/architecting-for-the-aws-cloud-best-practices/>

Design principles for Cloud Solutions

1. Design for failure
2. Loose coupling
3. Implement “Elasticity”
4. Build Security in every layer
5. Don't fear constraints
6. Think Parallel
7. Leverage different storage options

- **Avoid single points of failure**
- Use **queues to decouple** : *Job A passes the output to a queue for Job B to pull from*
- **Idempotent and Metadata driven** scaling - *new jobs/actors should be bootstrapped to ask “who am I, what am I supposed to do?”*
- In a network, location/perimeter security is useless - so wrap all the components with **ownership, audit.**
- Memory, IOPS distribution by **sharding, Caching** for performance
- **Divide, Parallelize and conquer.**
Serial/Sequential days are gone.
- **One database does not fit all** anymore. Think different stores for the outcome (Graph, Key-value, OLAP)

Learn from the industry leaders..



“To achieve good end-to-end latency, we wrote our ETL as a streaming job. By having it constantly running, we’re able to incrementally fill discrete hourly buckets as the data arrives. This gives us better latency compared to a batch job that exported data once at the end of every hour.”

- Igor Maravić, Spotify

*Watch on youtube our Cloud Next presentation they did talking about how they shutdown the largest Hadoop cluster in Europe



Spotify has a great published story about their migration to cloud.

BigQuery	10 million queries per month on scanning over 500PBs
Pub/Sub	Over 1 Trillion messages/day - under 400ms latency
Dataflow	5000+ dataflow jobs a day
Dataproc	Lots of dynamically spun up job scoped clusters

“The best way to avoid failure is to fail constantly”

NETFLIX MEDIA CENTER

11 February 2016

Completing the Netflix Cloud Migration



Our journey to the cloud at Netflix began in August of 2008, when we experienced a major database corruption and for three days could not ship DVDs to our members. That is when we realized that we had to move away from vertically scaled single points of failure, like relational databases in our datacenter, towards highly reliable, horizontally

Simian Army

- Kill/inspect running instances
 - Chaos Monkey
 - Janitor Monkey
 - Security Monkey
 - Conformity Monkey
 - Chaos Gorilla*
 - Chaos Kong*

<https://github.com/Netflix/SimianArmy>

https://github.com/Netflix/security_monkey

10/10/14

@SonOfGarr



WRITTEN BY



YURY IZRAILEVSKY
Vice President, Cloud and Platform Engineering

Thank You

END