DATA WRANGLING FOR ETL ENTHUSIASTS

MOHAMED KABIRUDDIN
CLOUD SOLUTIONS ARCHITECT (DATA)
MICROSOFT

ETL PROCESS

LANDING

STAGING

DIMENSIONAL MODEL

DID I JUST SPEND 10 HOURS PERFECTING THAT LOOKUP?

AND STILL APPLY INDEXES TO GAIN PERFORMANCE?

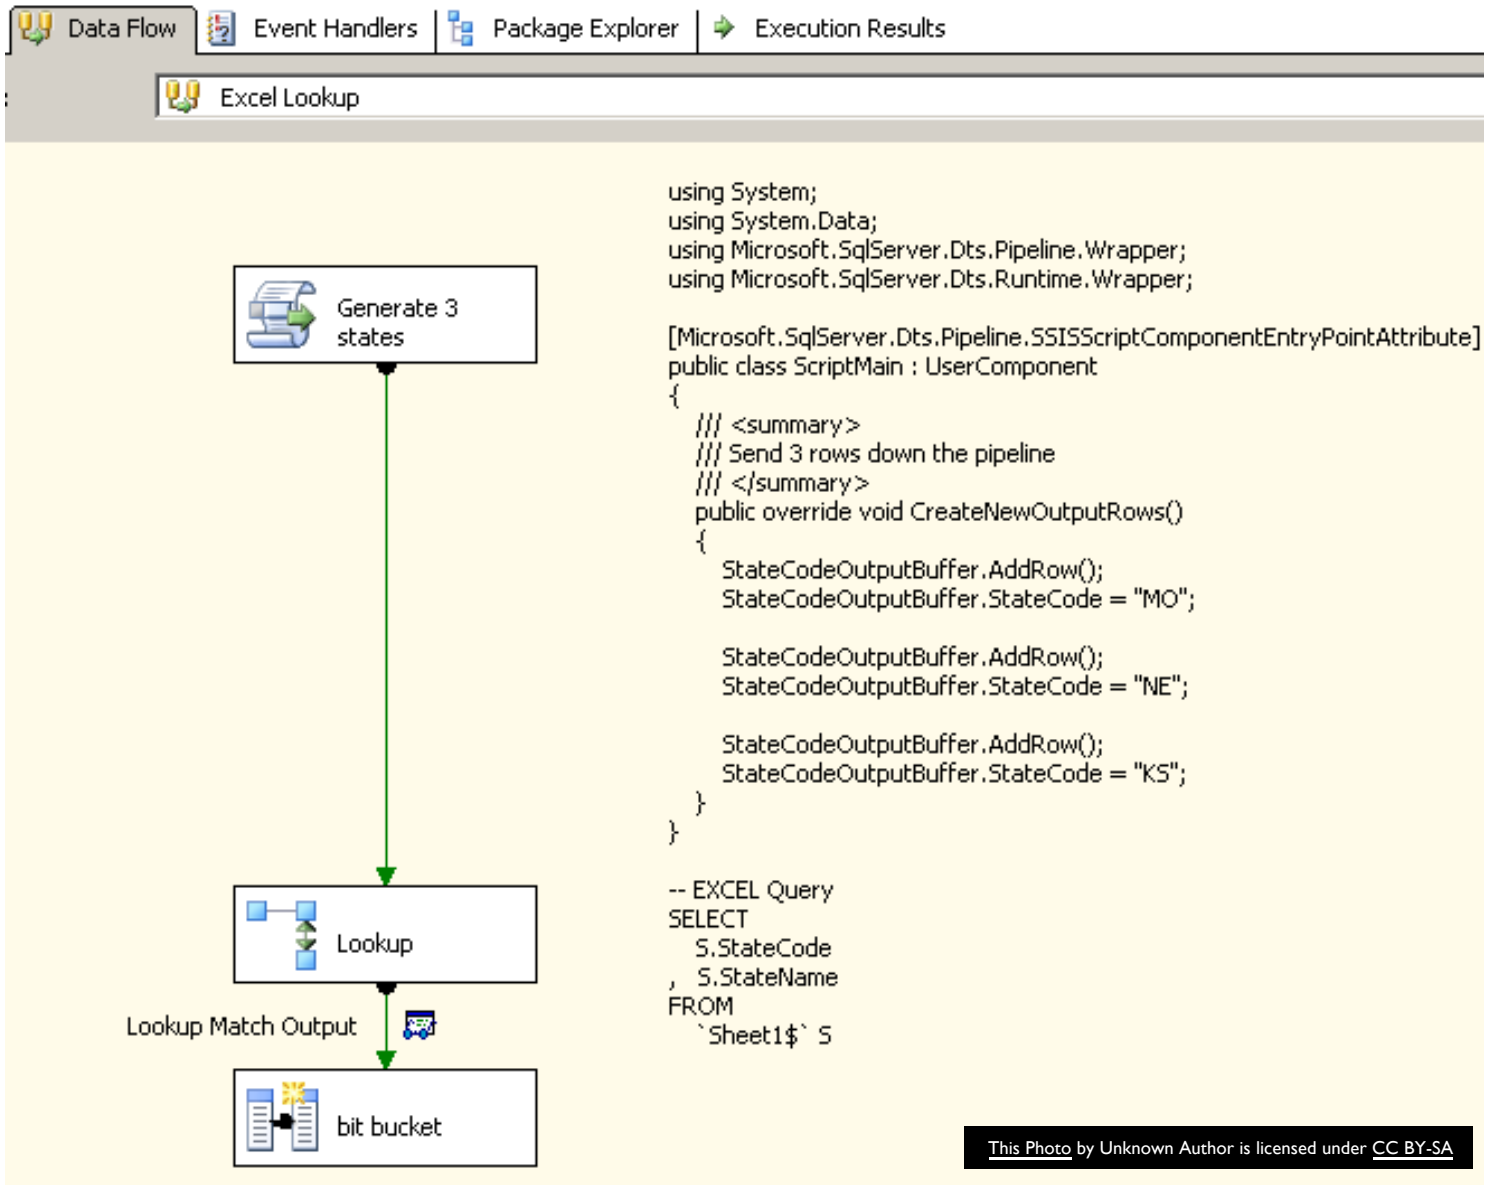

CONSTRUCT THAT POWERFUL SURROGATE KEY

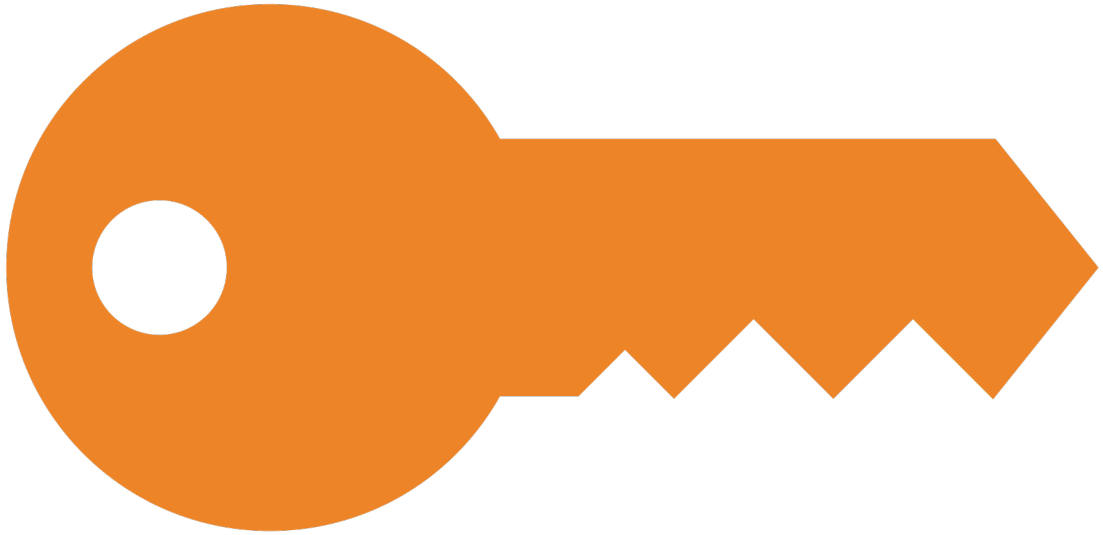

ULATIMATE GOAL:

KEEP THE
DATA WAREHOUSE

1. UPDATED
2. RELEVANT
3. OPERATIONAL

DATA STAGES IN WRANGLING

RAW

REFINED

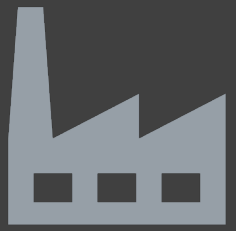

PRODUCTION

DATA LAKE DESIGN CONSIDERATIONS

Data Lake Zones

Transient Landing Zone

Temporary storage of data to meet regulatory and quality control requirements. Limited access. May not be required depending on requirements.

Raw Zone

Original source of data ready for consumption. Metadata publicly available but access to data still limited.

Trusted Zone

Standardized and enriched datasets ready for consumption to those with appropriate role-based access. Metadata available to all.

Curated/Refined Zone

Data transformed from Trusted Zone to meet specific business requirements.

Sandbox Zone

Playground for Data Scientists for ad hoc exploratory use cases.

Data Governance Considerations

Security and Compliance

Access Control

Encryption

Row-Level Security

Metadata Management

Data Quality

Metadata Management

Lifecycle Management

ARE DATA WRANGLING AND ETL THE SAME THEN?

- Data wrangling is the process involved in transforming or preparing data for analysis
- Consider ETL to be one type of data wrangling, specifically a type of data wrangling managed and overseen by an organization's shared services or IT organization.
- Data wrangling can also be handled by business users in desktop tools like Excel, or by data scientists in coding languages like Python or R

Principles of Data Wrangling

by Connor Carreras; Jeffrey Heer; Sean Kandel; Joseph M. Hellerstein; Tye Rattenbury

Published by O'Reilly Media, Inc., 2017

O'REILLY

TOOLSET FROM MICROSOFT

- Power BI
- Excel
- SSIS
- T-SQL
- U-SQL
- Polybase
- Azure Data Explorer
- Azure Data Factory
- Azure Stream Analytics
- Azure HD Insight (R & Python)
- Azure Databricks (R, Python and SparkSQL)

MODERN DATA WAREHOUSING

Canonical operations

Load and ingest

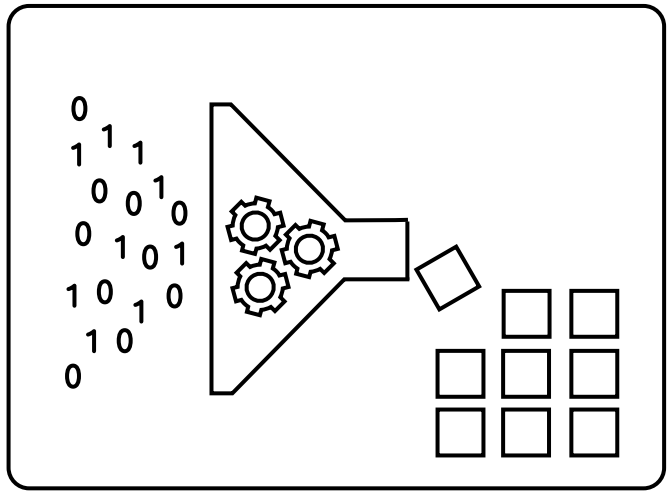

Transfer and store

Process

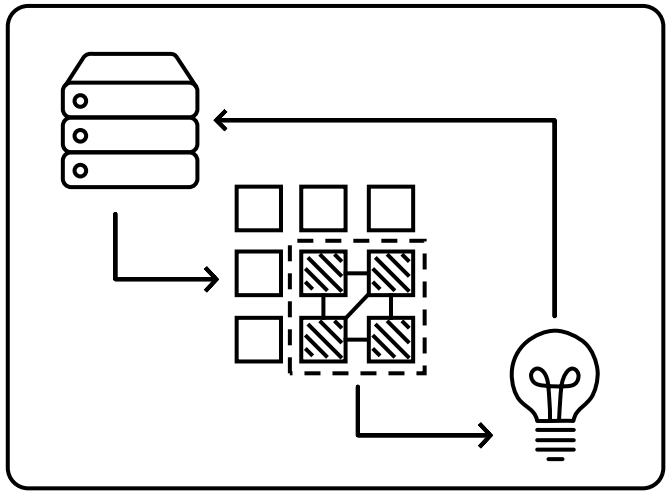

Process and clean

Serve

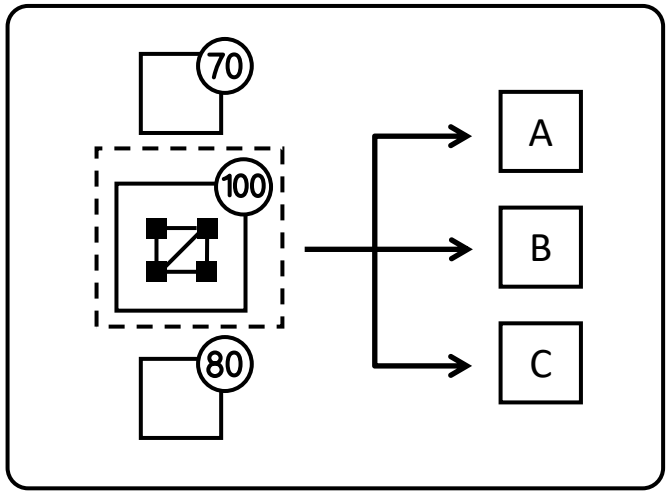

Serve and analyze

MODERN DATA WAREHOUSING PATTERN IN AZURE

Data processing with Azure Databricks

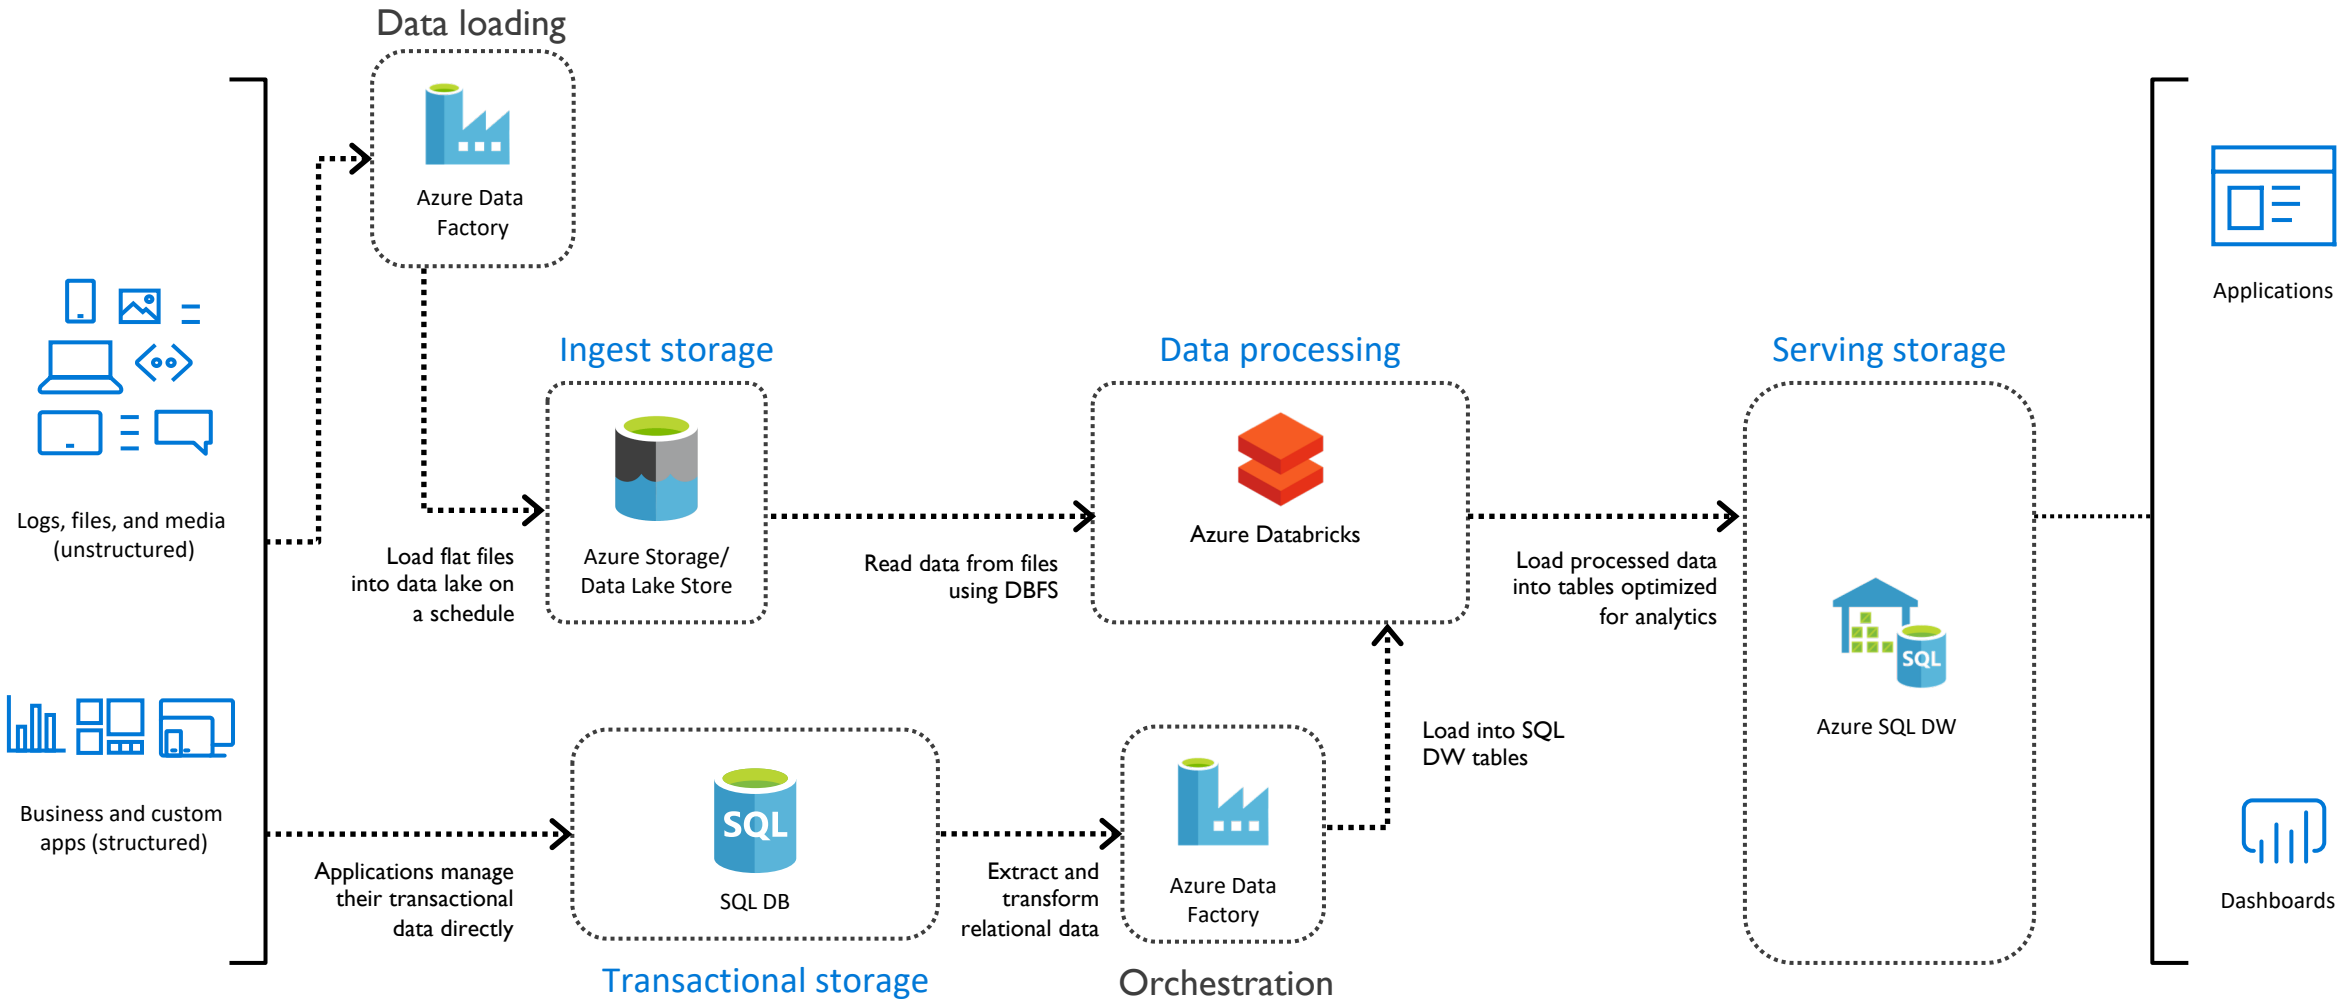

New Data Flow

Mapping Data Flow

Code free data transformation at scale

Wrangling Data Flow (Preview)

Code free data preparation at scale

DATA FACTORY - DATAFLOWS

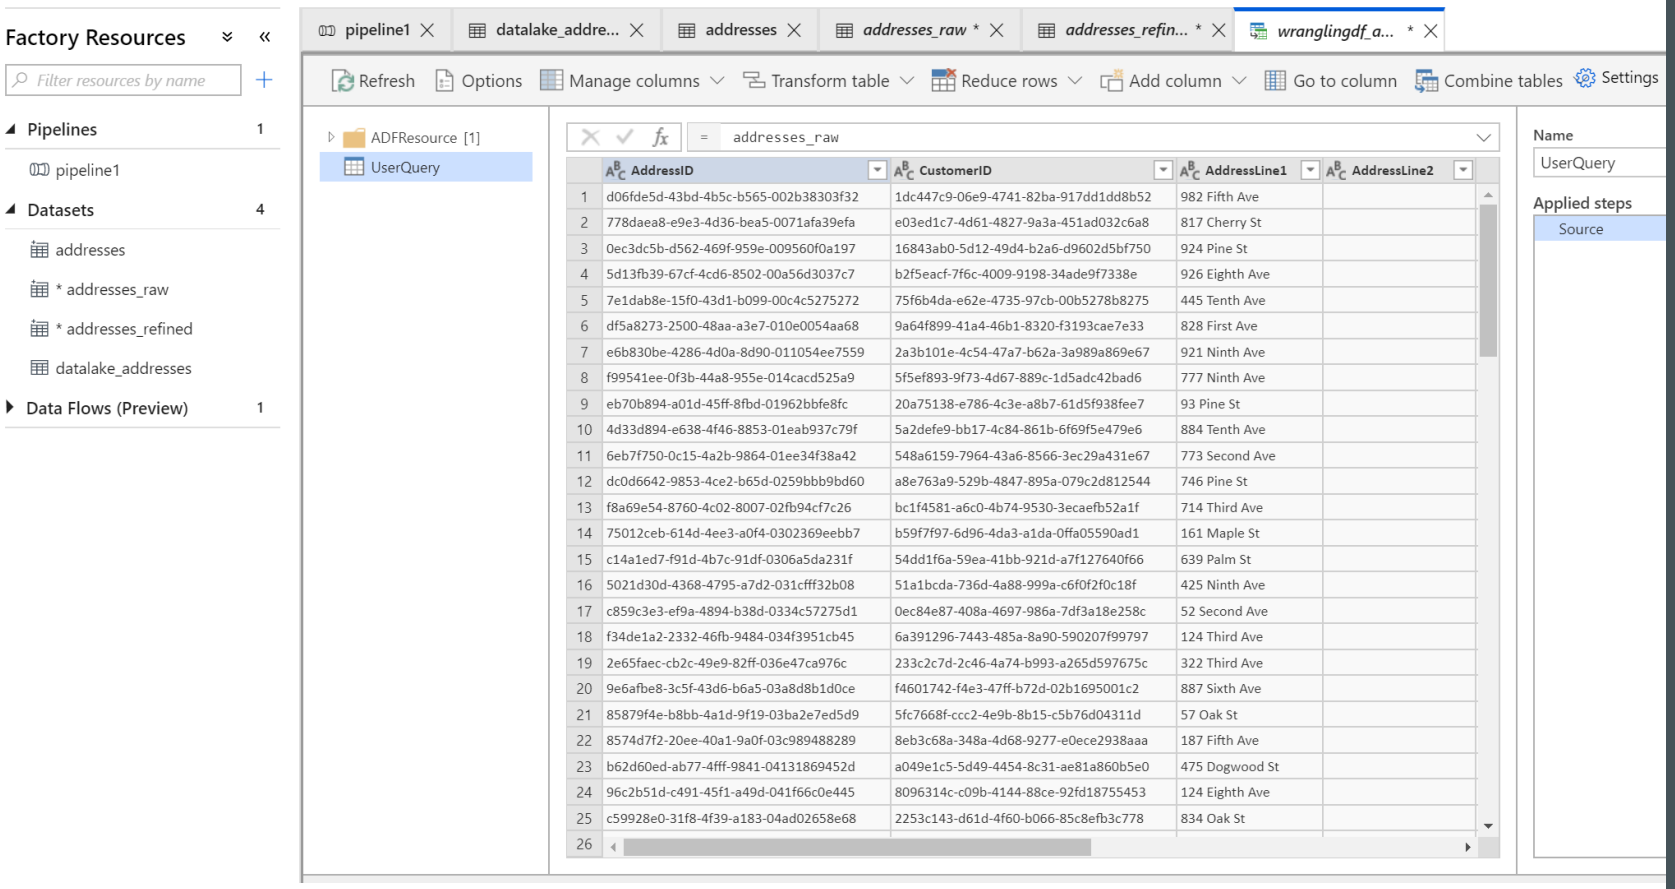

WRANGLING DATAFLOWS

MAPPING DATA FLOW

No-code data transformation @ scale

Data cleansing, transformation, aggregation, conversion, etc.

Cloud scale via Spark execution

Easily build resilient data flows

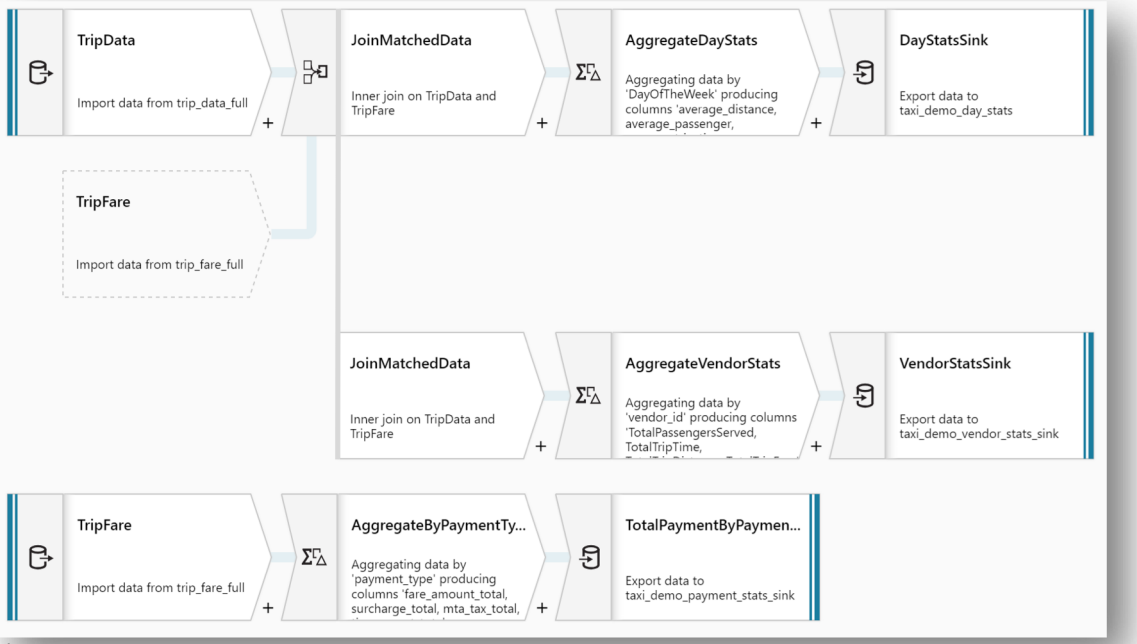

... not

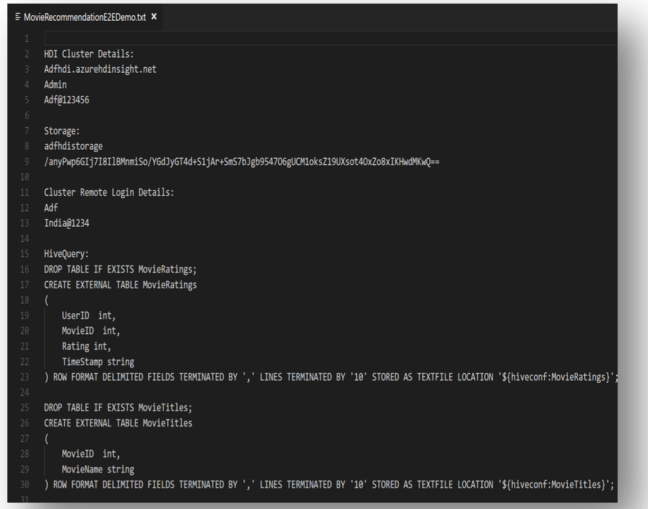

AZURE DATABRICKS

A fast, easy and collaborative Apache® Spark™ based analytics platform optimized for Azure

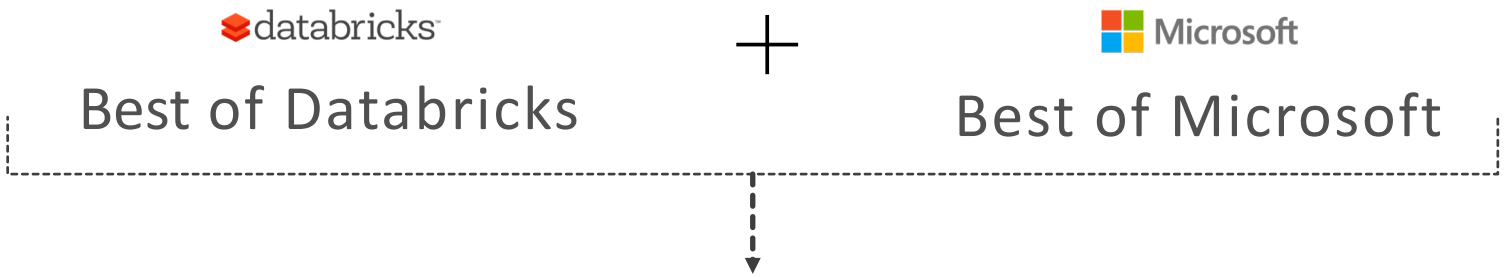

- Designed in collaboration with the founders of Apache Spark
- One-click set up; streamlined workflows
- Interactive workspace that enables collaboration between data scientists, data engineers, and business analysts.
- Native integration with Azure services (Power BI, SQL DW, Cosmos DB, ADLS, Azure Storage, Azure Data Factory, Azure AD, Event Hub, IoT Hub, HDInsight Kafka, SQL DB)
- Enterprise-grade Azure security (Active Directory integration, compliance, enterprise-grade SLAs)

AZURE DATABRICKS NOTEBOOKS

Notebooks are a popular way to develop, and run, Spark Applications

Notebooks are not only for authoring Spark applications but can be *run/executed directly* on clusters

- Shift+Enter
- click the ▶ at the top right of the cell in a notebook
- Submit via Job

Fine grained permissions support so they can be *securely shared* with colleagues for collaboration

Notebooks are well-suited for prototyping, rapid development, exploration, discovery and iterative development

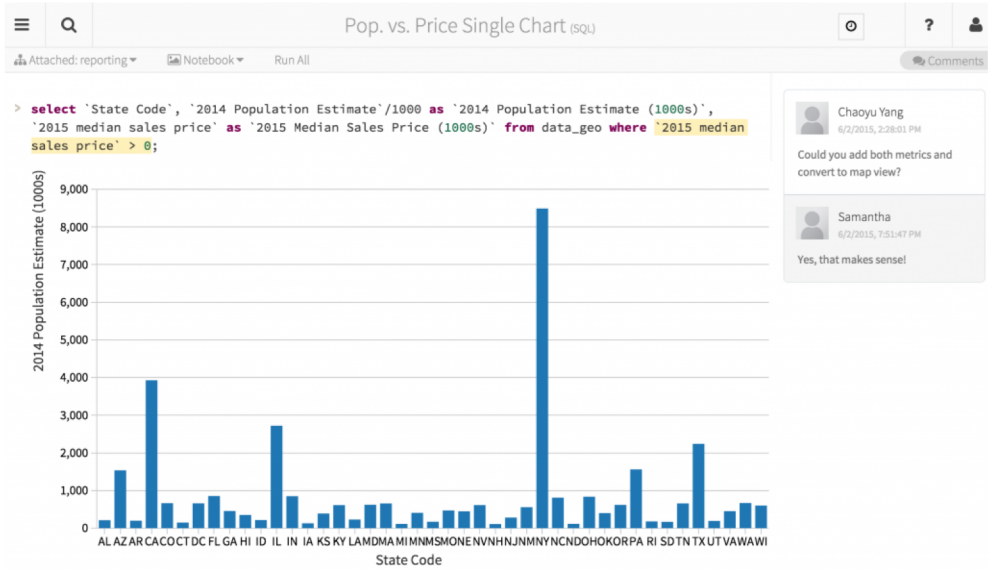

With Azure Databricks notebooks you have a default language but you can mix multiple languages in the same notebook:

`%python`	Allows you to execute python code in a notebook (even if that notebook is not python)
`%sql`	Allows you to execute sql code in a notebook (even if that notebook is not sql).
`%r`	Allows you to execute r code in a notebook (even if that notebook is not r).
`%scala`	Allows you to execute scala code in a notebook (even if that notebook is not scala).
`%sh`	Allows you to execute shell code in your notebook.
`%fs`	Allows you to use Databricks Utilities - dbutils filesystem commands.
`%md`	To include rendered markdown