

Hi, I'm Claire 🖐️✍️

@clairebcarroll



**FISHTOWN**  
ANALYTICS



# Data engineers shouldn't write DDL

Claire Carroll, dbt Community Manager

# What is DDL?

- Data definition language
- Queries that define or alter the objects in a database
- Any query other than **select**:
  - **create, alter, drop**
  - **truncate, insert** and **update**\*

```
create schema sandpit_claire;
```

```
create table sandpit_claire.stg_payments as (  
  select  
    id as payment_id,  
    order_id,  
    payment_method,  
  
    -- amount is currently stored in cents  
    amount / 100 as amount  
  
  from raw_jaffle_shop.payments  
);
```

# How to write DDL

**select**

id **as** payment\_id,  
order\_id,  
payment\_method,

-- amount is currently stored in cents  
amount / 100 **as** amount

**from** raw\_jaffle\_shop.payments;

```
create table sandpit_claire.stg_payments as (  
  select  
    id as payment_id,  
    order_id,  
    payment_method,  
  
    -- amount is currently stored in cents  
    amount / 100 as amount  
  
  from raw_jaffle_shop.payments  
);
```



```
create schema sandpit_claire;
```

```
create table sandpit_claire.stg_payments as (  
  select
```

```
    id as payment_id,  
    order_id,  
    payment_method,
```

```
    -- amount is currently stored in cents  
    amount / 100 as amount
```

```
from raw_jaffle_shop.payments  
);
```

```
create schema if not exists sandpit_claire;
```

```
drop table sandpit_claire.stg_payments;
```

```
create table sandpit_claire.stg_payments as (  
  select
```

```
    id as payment_id,  
    order_id,  
    payment_method,
```

```
    -- amount is currently stored in cents  
    amount / 100 as amount
```

```
from raw_jaffle_shop.payments  
);
```

```
create schema if not exists sandpit_claire;
```

```
drop table if exists sandpit_claire.stg_payments;
```

```
create table sandpit_claire.stg_payments as (
```

```
  select
```

```
    id as payment_id,
```

```
    order_id,
```

```
    payment_method,
```

```
    -- amount is currently stored in cents
```

```
    amount / 100 as amount
```

```
from raw_jaffle_shop.payments
```

```
);
```

```
create schema if not exists sandpit_claire;  
  
begin;  
  
drop table if exists sandpit_claire.stg_payments;  
  
create table sandpit_claire.stg_payments as (  
  select  
    id as payment_id,  
    order_id,  
    payment_method,  
  
    -- amount is currently stored in cents  
    amount / 100 as amount  
  
  from raw_jaffle_shop.payments  
);  
  
commit;
```

```
create schema if not exists sandpit_claire;

begin;

drop table if exists sandpit_claire.stg_payments;

create table sandpit_claire.stg_payments as (
  select
    id as payment_id,
    order_id,
    payment_method,

    -- amount is currently stored in cents
    amount / 100 as amount

  from raw_jaffle_shop.payments
);

commit;
```

```
create view sandpit_claire.payments_agg as (
  select
    order_id,

    count(*) as n_payments,
    sum(
      case
        when payment_type = 'credit_card'
        then amount
        else 0
      end
    ) as credit_card_amount

  from sandpit_claire.stg_payments

  group by 1
);
```

```
create schema if not exists sandpit_claire;  
  
begin;  
  
drop table if exists sandpit_claire.stg_payments;  
  
create table sandpit_claire.stg_payments as (  
  select  
    id as payment_id,  
    order_id,  
    payment_method,  
  
    -- amount is currently stored in cents  
    amount / 100 as amount  
  
  from raw_jaffle_shop.payments  
);  
  
commit;
```

```
create view sandpit_claire.payments_agg as (  
  select  
    order_id,  
  
    count(*) as n_payments,  
    sum(  
      case  
        when payment_type = 'credit_card'  
        then amount  
        else 0  
      end  
    ) as credit_card_amount  
  
  from sandpit_claire.stg_payments  
  
  group by 1  
) with no schema binding;
```

# DDL optimizations

- Order of execution
- Atomic transactions
- Parameterized schemas
- Performance optimization
- Incremental builds
- Introspecting the warehouse

**Good DDL is complex**



# Technical challenges

- Harder to do performance optimization
- Harder to switch platforms
- Any features on top of this require a lot of work:
  - Testing
  - Automated documentation

# Organizational challenges

- Hard for others to understand

```
create schema if not exists sandpit_claire;  
  
begin;  
  
drop table if exists sandpit_claire.stg_payments;  
  
create table sandpit_claire.stg_payments as (  
  select  
    id as payment_id,  
    order_id,  
    payment_method,  
  
    -- amount is currently stored in cents  
    amount / 100 as amount  
  
  from raw_jaffle_shop.payments  
);  
  
commit;
```



```
create view sandpit_claire.payments_agg as (  
  select  
    order_id,  
  
    count(*) as n_payments,  
    sum(  
      case  
        when payment_type = 'credit_card'  
        then amount  
        else 0  
      end  
    ) as credit_card_amount  
  
  from sandpit_claire.stg_payments  
  
  group by 1  
) with no schema binding;
```

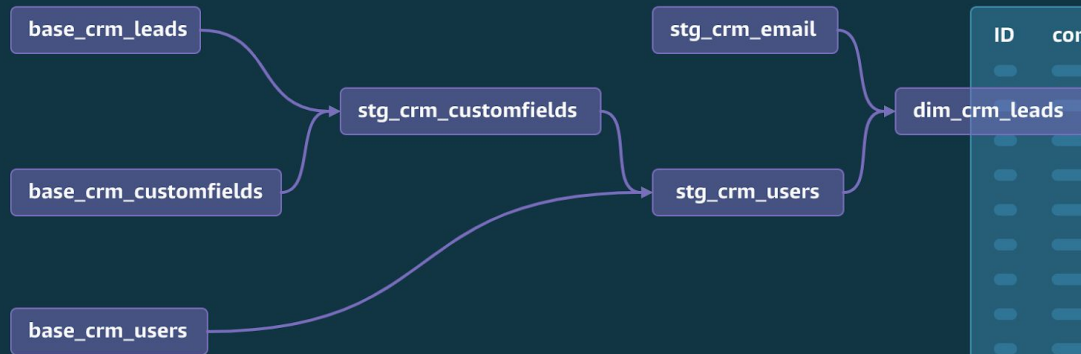
# Organizational

- Hard for others to understand
- Become a bottleneck
- Low job satisfaction

**How to not write DDL**

# Use a framework

- The framework generates the boilerplate code
- Options:
  - dbt ()
  - Roll your own ()
  - General DE platform: Airflow / Dagster / Prefect / Luigi
- Further reading: [Maintainable ETLs](#) – StitchFix



ID	company	lead_email	lead_owner	status	c...

```
SELECT
  leads.id,
  leads.company,
  leads.email,
  owners.name as lead_owner,
  leads.status

FROM {{ ref('stg_crm_leads') }} AS leads

LEFT JOIN {{ ref('stg_crm_salespeople') }} AS owners
  ON leads.owner_id = owners.id
```

# Benefits of a framework

## **Technological**

- Most problems have already been thought about
- Good frameworks are extensible
- Allow you to focus on core business logic

## **Organizations**

- Analysts are empowered to own the transformation layer
- Emergence of the analytics engineer



Overview - DBT Docs

127.0.0.1:8080/#!model/model.jaffle\_shop.fct\_orders

dbt\_

Search for models...

Overview

ProjectDatabase

Sources

- raw\_jaffle\_shop

Models

- jaffle\_shop
  - models
    - marts
      - core
        - dim\_customers
        - fct\_orders
        - intermediate
      - staging

fct\_orders table

DetailsDescriptionColumnsSQL

Details

TAGS	OWNER	TYPE	PACKAGE	RELATION	ADVANCED
untagged	claire	table	jaffle_shop	claire.jaffle_shop.dev.fct_orders	Show

Description

This table has basic information about orders, as well as some derived facts based on payments

Columns

COLUMN	TYPE	DESCRIPTION	TESTS	MORE?
order_id	integer	This is a unique identifier for an ...	U N	>
customer_id	integer	Foreign key to the customers table	N F	>
order_date	date	Date (UTC) that the order was pl...		>
status	character varying	Orders can be one of the followin...		>
credit_card_amount	bigint	Amount of the order (AUD) paid f...	N	>
coupon_amount	bigint	Amount of the order (AUD) paid f...	N	>
bank_transfer_amount	bigint	Amount of the order (AUD) paid f...	N	>
gift_card_amount	bigint	Amount of the order (AUD) paid f...	N	>

# Benefits of a framework

## **Technological**

- Most problems have already been thought about
- Good frameworks are extensible
- Allow you to focus on core business logic

## **Organizational**

- Analysts are empowered to own the transformation layer
- Emergence of the analytics engineer

Data engineers shouldn't write DDL, but...

**Data engineers should**  
*deeply understand* DDL

# Concurrent trends

Where possible:

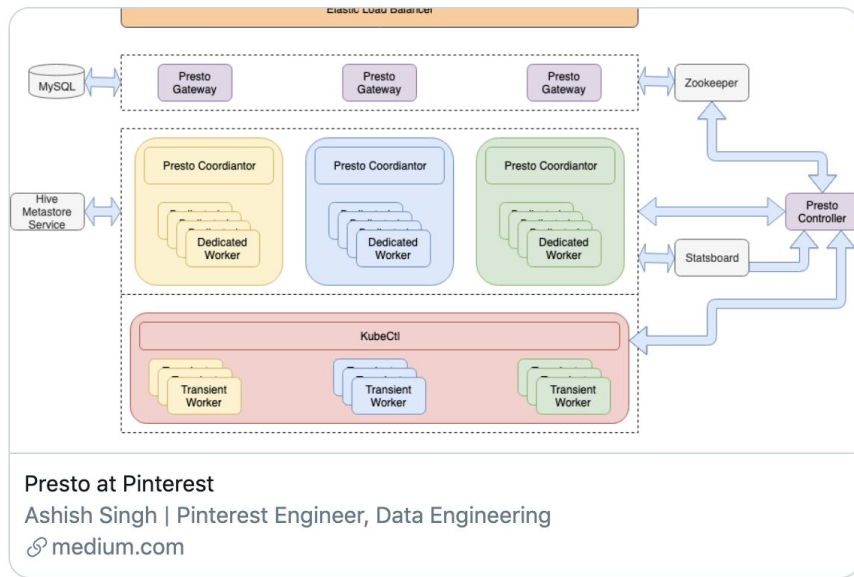
- Data engineers shouldn't write data ingestion
- Data engineers shouldn't write data collectors
- Data engineers shouldn't build data warehouse technology

chrismunns  
@chrismunns



So Pinterest built a bad ass Presto cluster system across thousands of cores and TBs of data:

[medium.com/@Pinterest\\_Eng...](https://medium.com/@Pinterest_Engineering/100-days-of-presto-at-pinterest-4a0000000000) but YOU should probably just use Athena & RedShift unless you've got an entire TEAM to dedicate to big data systems..



So what *do* data  
engineers do?

# Extending frameworks

- Custom materializations in dbt
- Building features for dbt

# Data warehouse performance

- Performant data model design
- Performance tuning



# Deployment

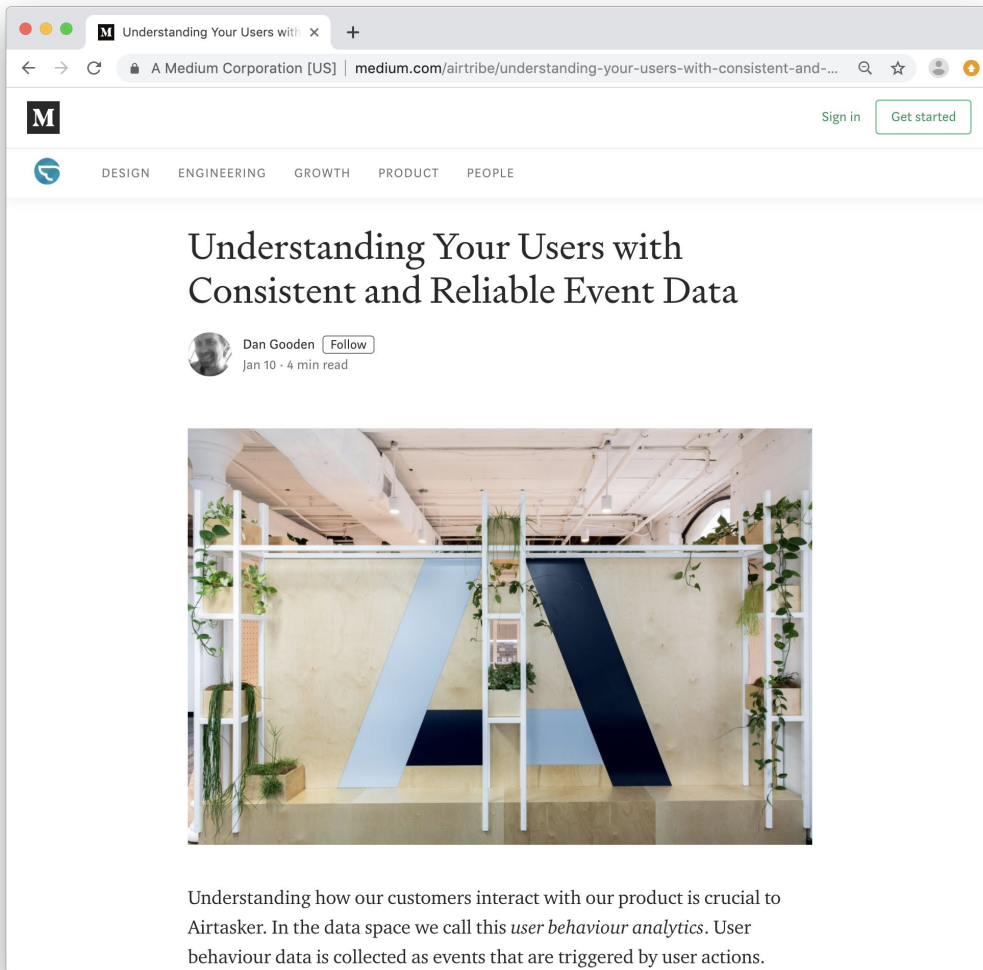
- CI
- CD
- Orchestration
- Credentials storage
- Monitoring
- Alerting

# Custom data ingestion

- For APIs that aren't on Stitch / Singer
- Writing Singer taps

# Source data design

- Working with core product team
- Data quality enforcement



# Data platform infrastructure

- Luigi to Dagster?
- Redshift to Snowflake?

# Working with analysts

- Teach them engineering concepts
- Learn how to make their lives more productive

# The role of data engineers

- Extend frameworks
- Data warehouse performance
- Deployment
- Custom data ingestion
- Source data design and data quality enforcement
- Data platform infrastructure
- Working with analysts

# Thanks!

**Any questions?**

- @getdbt
- getdbt.com
- learn.getdbt.com