

Turning the wheel - Streaming 4.4 billion events with Apache Kafka

Streaming customer, policy and vehicle
information via Apache Kafka and MongoDB –
and what we learnt on the way ...

Simon Aubury

April 2019



Agenda

Tonight

- Context - what problem are we trying to solve
- Architecture of our data flow
- Kafka & Kafka Connect
- Challenges .. and solutions



We've a lot of data

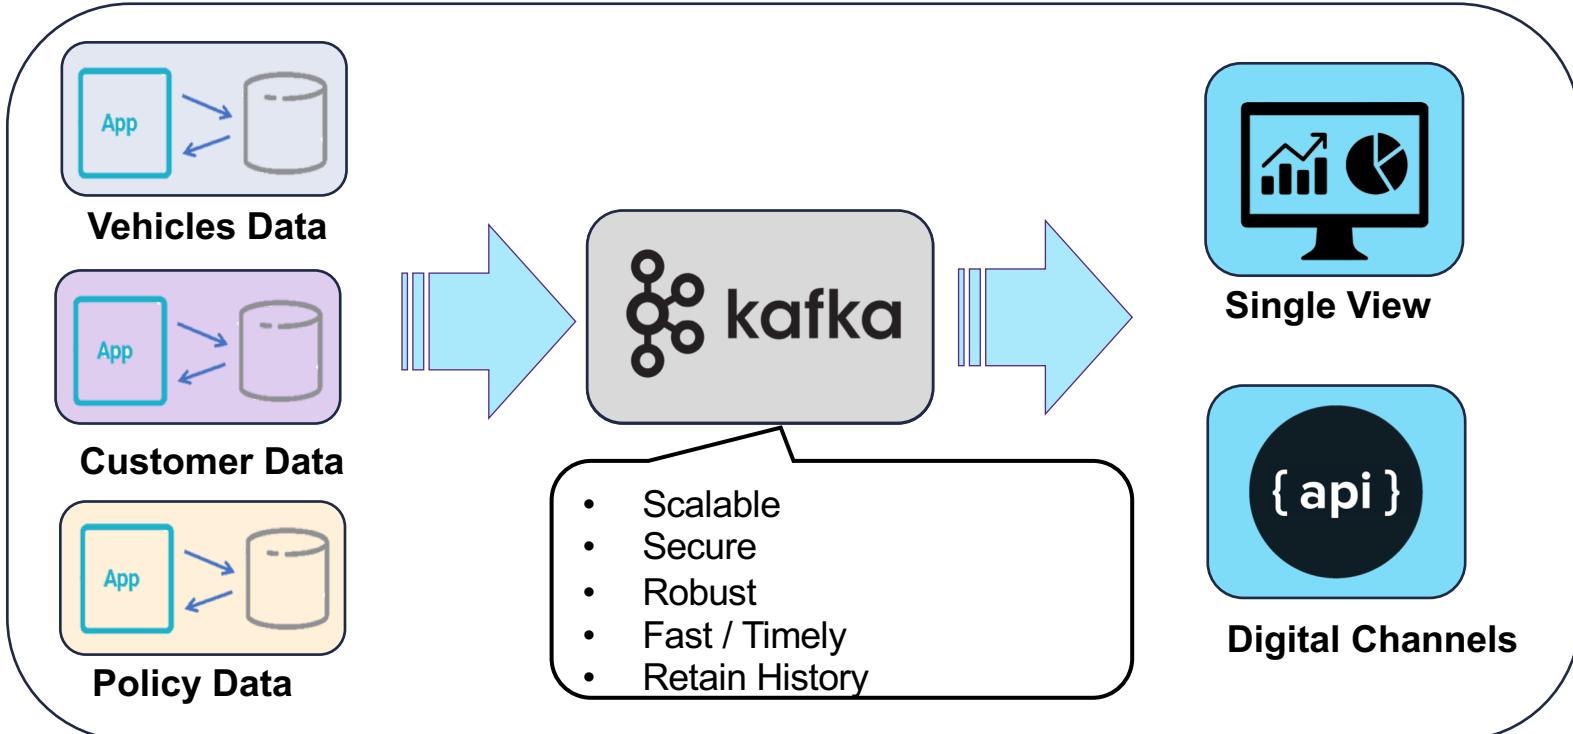
Context #1

mvyear	mvmake	mvmodel	mvbody
1886	RUDGE	PENNY FARTHING	MBIKE
1896	FORD	QUADRICYCLE	CONVT
1896	FORD	QUADRICYCLE	CONVT
1896	FORD	QUADRICYCLE	CONVT
1896	FORD	QUADRICYCLE	CONVT
1896	FORD	QUADRICYCLE	CONVT
1896	FORD	QUADRICYCLE	CONVT
1896	FORD	QUADRICYCLE	CONVT
1896	FORD	QUADRICYCLE	CONVT
1900	MERCEDES		CONVT



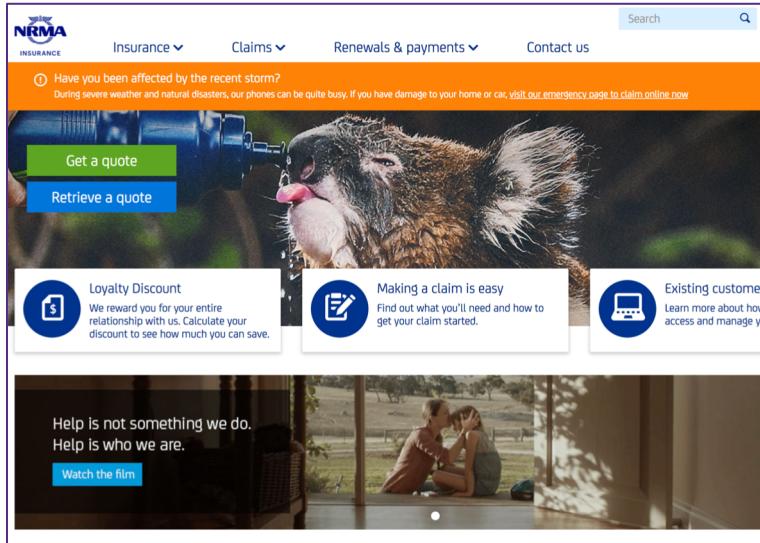
We've a lot of systems

Context #2

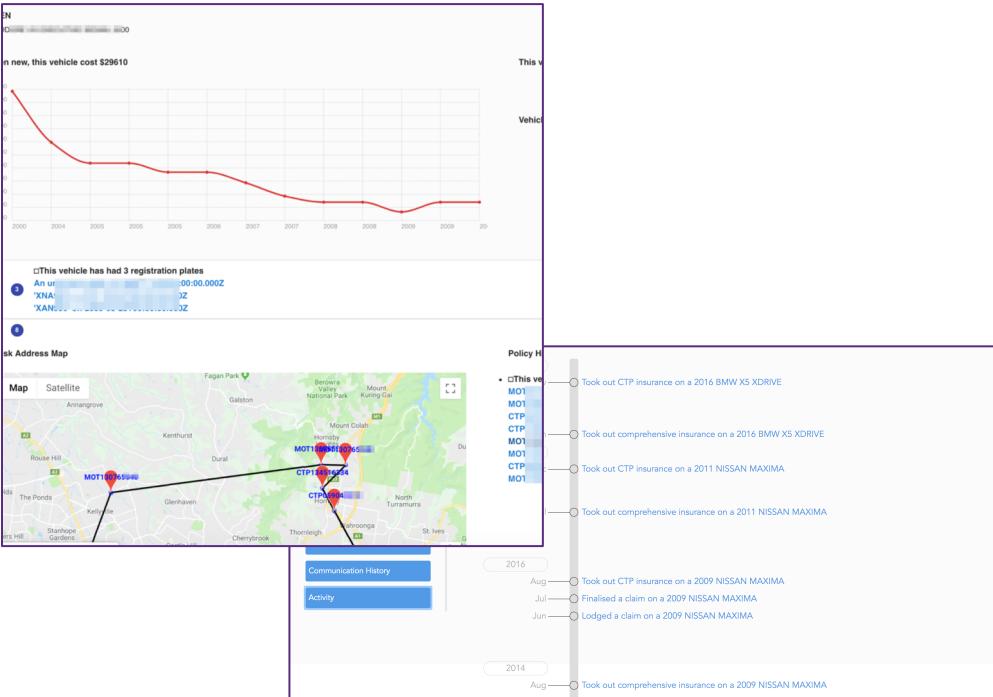


We want to tie it together

Context #3



The NRMA Insurance website features a prominent koala drinking from a water bottle on the homepage. The navigation bar includes links for Insurance, Claims, Renewals & payments, and Contact us. A search bar is at the top right. A message at the top left asks if users have been affected by the recent storm. Below the message are three main service buttons: Get a quote, Retrieve a quote, Loyalty Discount, Making a claim is easy, and Existing customer. A large video player at the bottom shows two people in a garden setting.



This screenshot displays a vehicle history report. At the top, a graph shows the value of a vehicle from 2003 to 2010, starting at approximately \$29,610 and fluctuating slightly. Below the graph, a message states: "This vehicle has had 3 registration plates". A map of Australia highlights several locations with blue dots, each labeled with a policy number: MOT12345678, CTP12345678, and CTP12345678. To the right, a sidebar lists policy history with entries for 2016, 2015, 2014, and 2013, detailing various insurance types like CTP, MOT, and comprehensive insurance, along with specific dates and descriptions of claims or policies taken out.

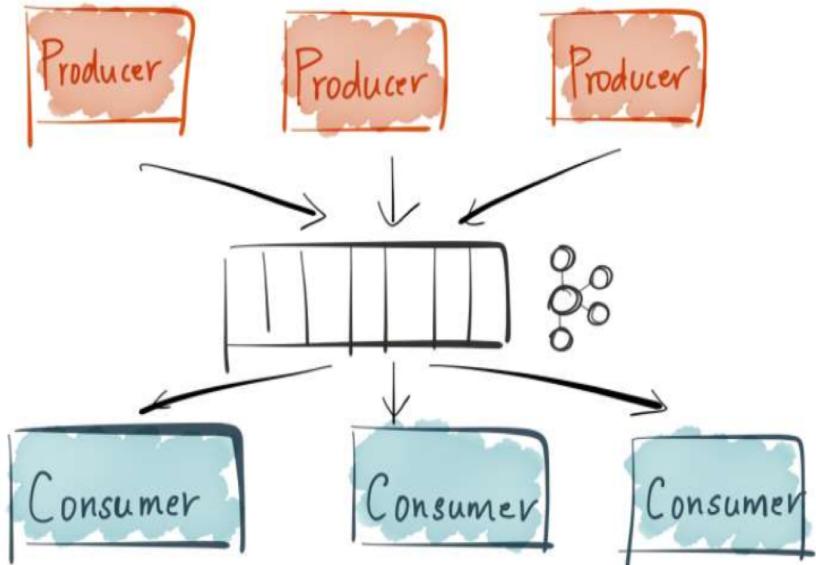


What is Kafka?

A very quick intro

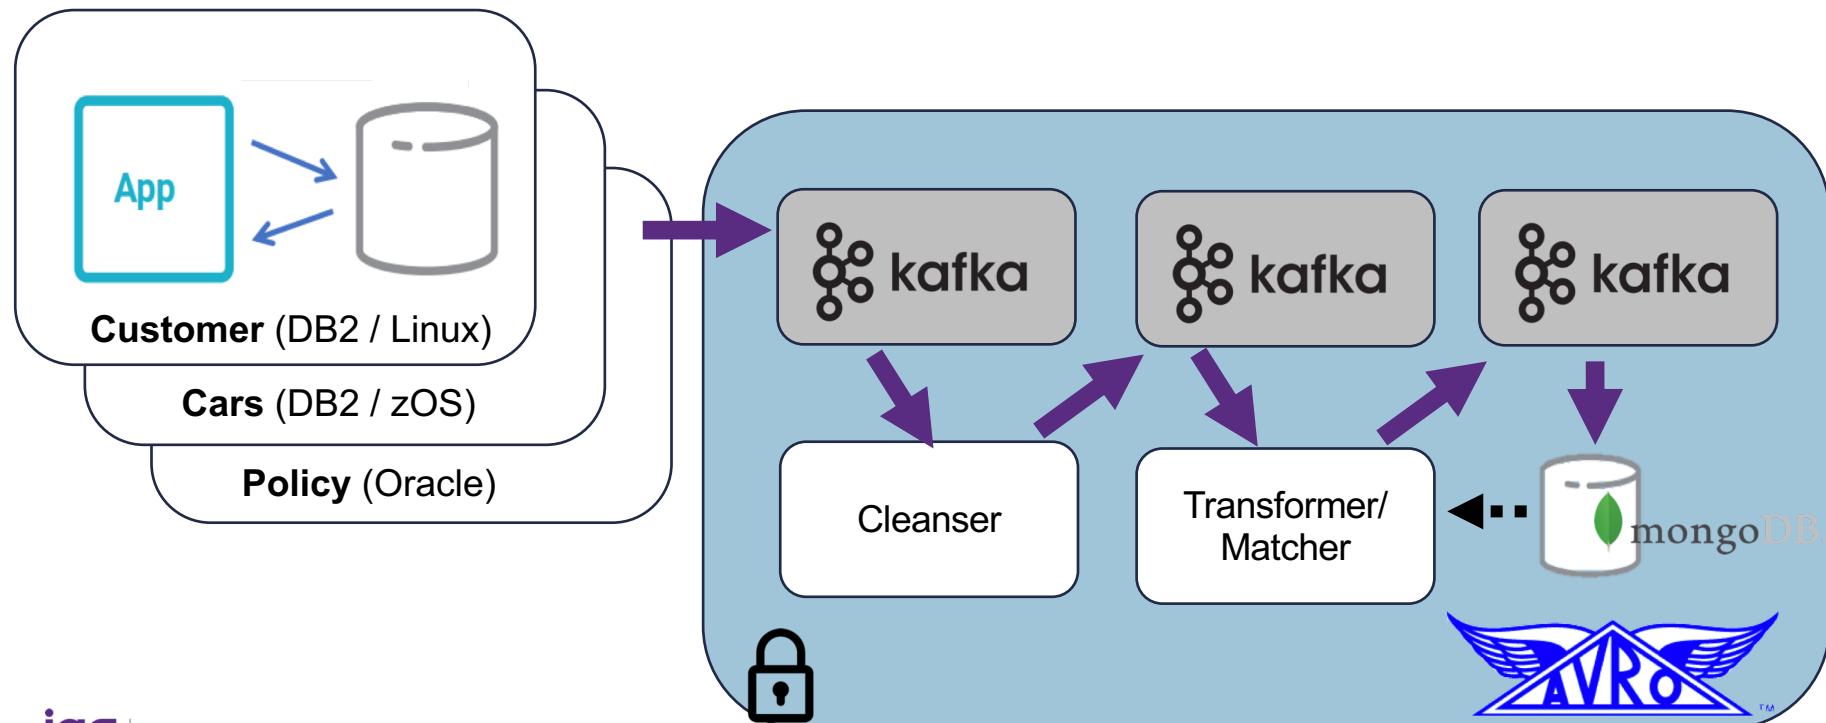
Apache Kafka : *Unified, high-throughput, low-latency platform for handling real-time data feeds*

- Originally developed by LinkedIn, open sourced in early 2011
- “The global commit log thingy”
- Kafka maintains feeds of messages in topics
- Appends ; ordered, immutable sequence



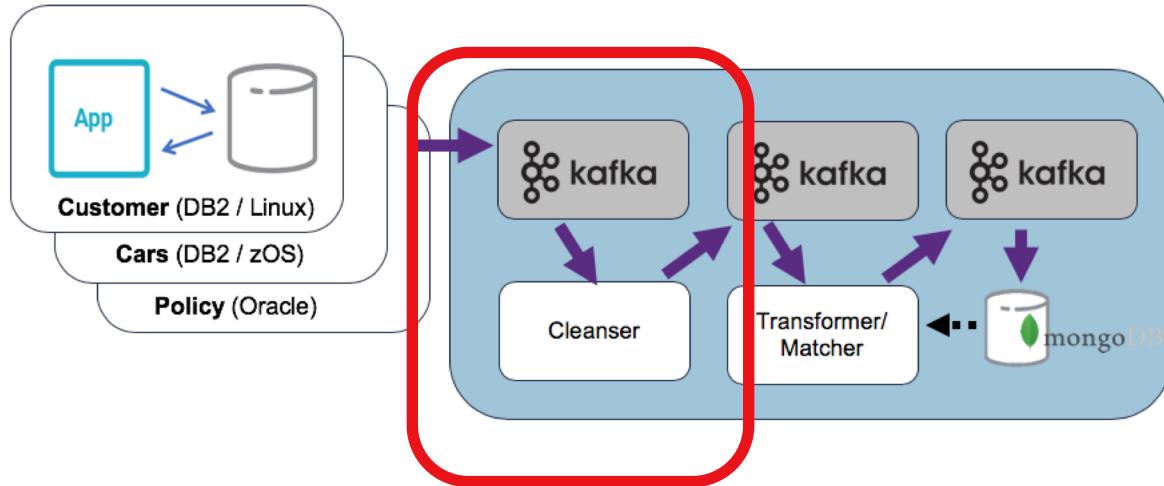
Architecture of our data flow

Lots of boxes



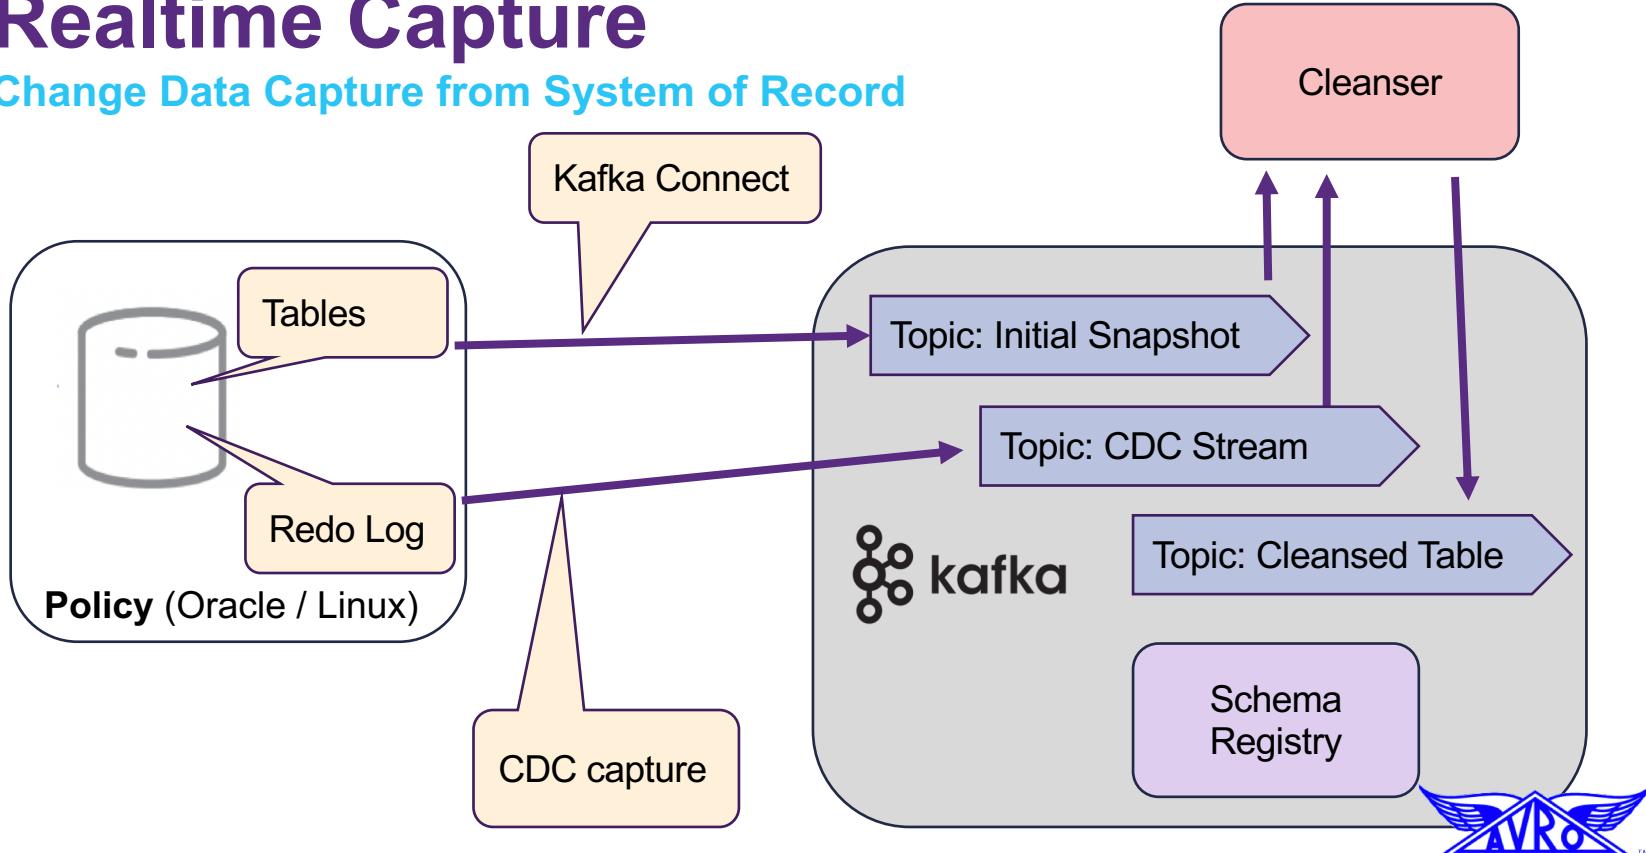
Part 1 - Extract

Source System Low Touch Data Acquisition



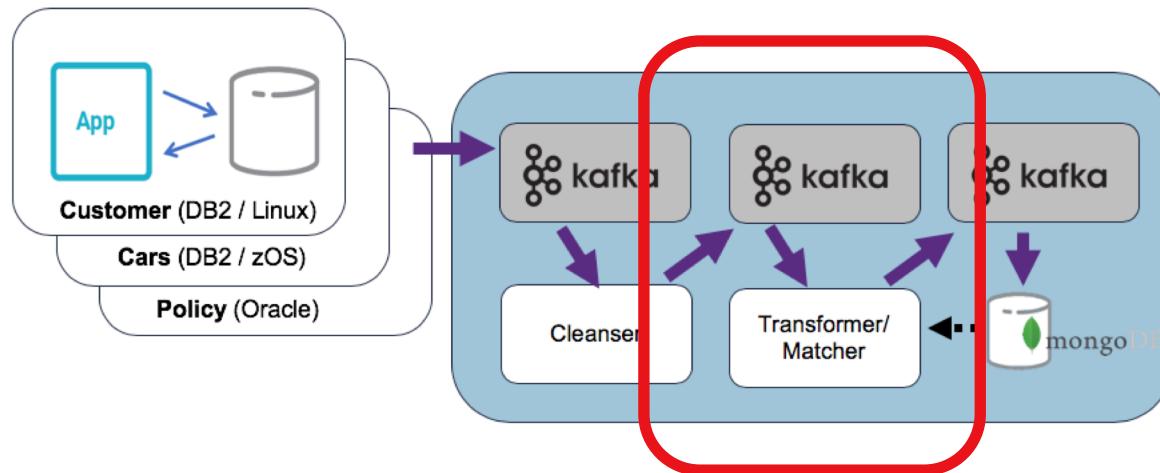
Realtime Capture

Change Data Capture from System of Record



Part 2 – Transformation & Matching

Finding stuff

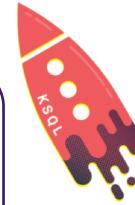


Transform & Match



Transformer/
Matcher

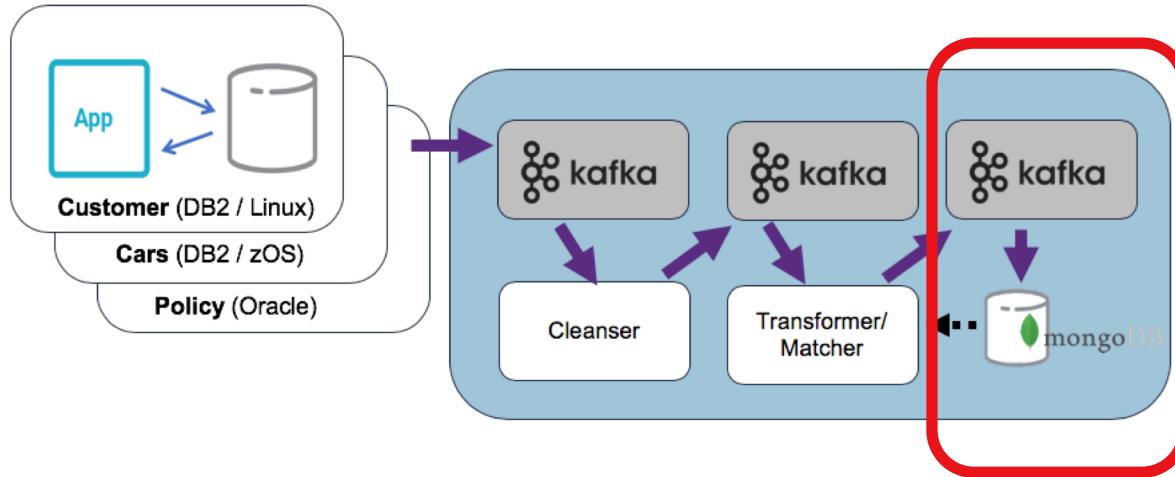
```
CREATE STREAM insurance_event_with_repairer AS \
SELECT *, geo_distance(iel.pc_lat, iel.pc_long, rct.lat, rct.long,
'km') AS dist_to_repairer_km
FROM insurance_event_with_location iel \
INNER JOIN repair_center_tab rct ON iel.pc_state = rct.repair_state;
```



```
final KTable<CcuserCcPolicyKey, ClaimDetails> claimDetailsTable = ccPiSorStreams.claim() KSt
.flatMapValues(statefulFilter(validClaim)) KStream<CcuserCcClaimKey, CcuserCcClaimEnvelope>
.mapValues((k, v) -> v.getAfter()) KStream<CcuserCcClaimKey, CcuserCcClaim>
.leftJoin(
    ccPiGlobalTables.brandExt(),
    (k, v) -> new BrandTypeCodeKey(v.getCLAIMNUMBER().substring(0, 3).toUpperCase()),
    (claim, brandExtEnv) -> {...}) KStream<CcuserCcClaimKey, ClaimDetails>
.groupByKey((k, v) -> new CcuserCcPolicyKey(v.getPolicyId()), serdes.groupedWith()) KGrouped
.reduce((prev, latest) -> latest, serdes.materializedAs( topicName: "claimDetails") KTabl
.join(
    policyTable,
    (claimDetails, policy) ->
        ClaimDetails.newBuilder(claimDetails)
            .setPolicyDetails(
                PolicyDetails.newBuilder()
                    .setBrand(policy.getDISTRIBUTOR())
                    .setPolicyNumber(policy.getPOLICYNUMBER())
                    .build())
            .build(),
    serdes.materializedAs( topicName: "claimDetailsWithPolicy"));
```

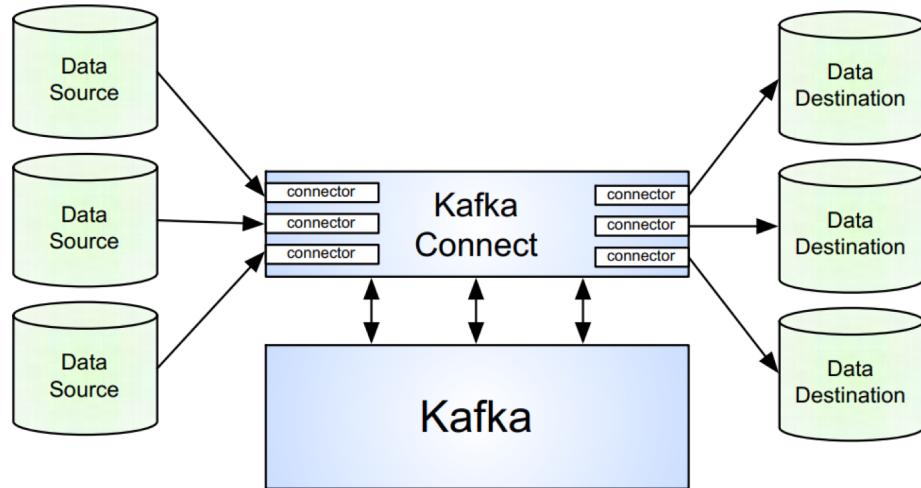
Part 3 – Serving Layer

Sinking Results



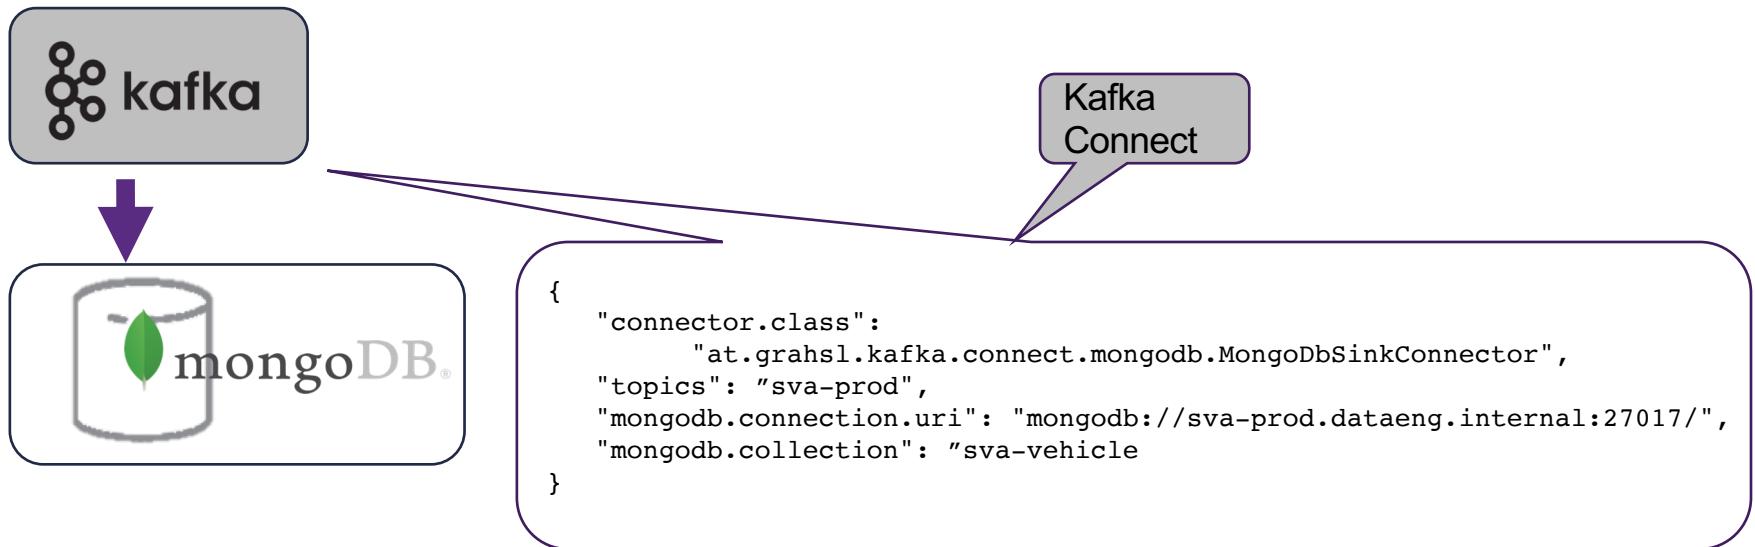
Kafka Connect

- Distributed, scalable, fault-tolerant service designed to reliably stream data between Kafka and other data systems
- **Source Connectors** import data from another system (e.g. a relational database into Kafka)
- **Sink Connectors** export data (e.g. the contents of a Kafka topic to an HDFS file).



Kafka Connect Sink

Writing to MongoDB Serving Layer



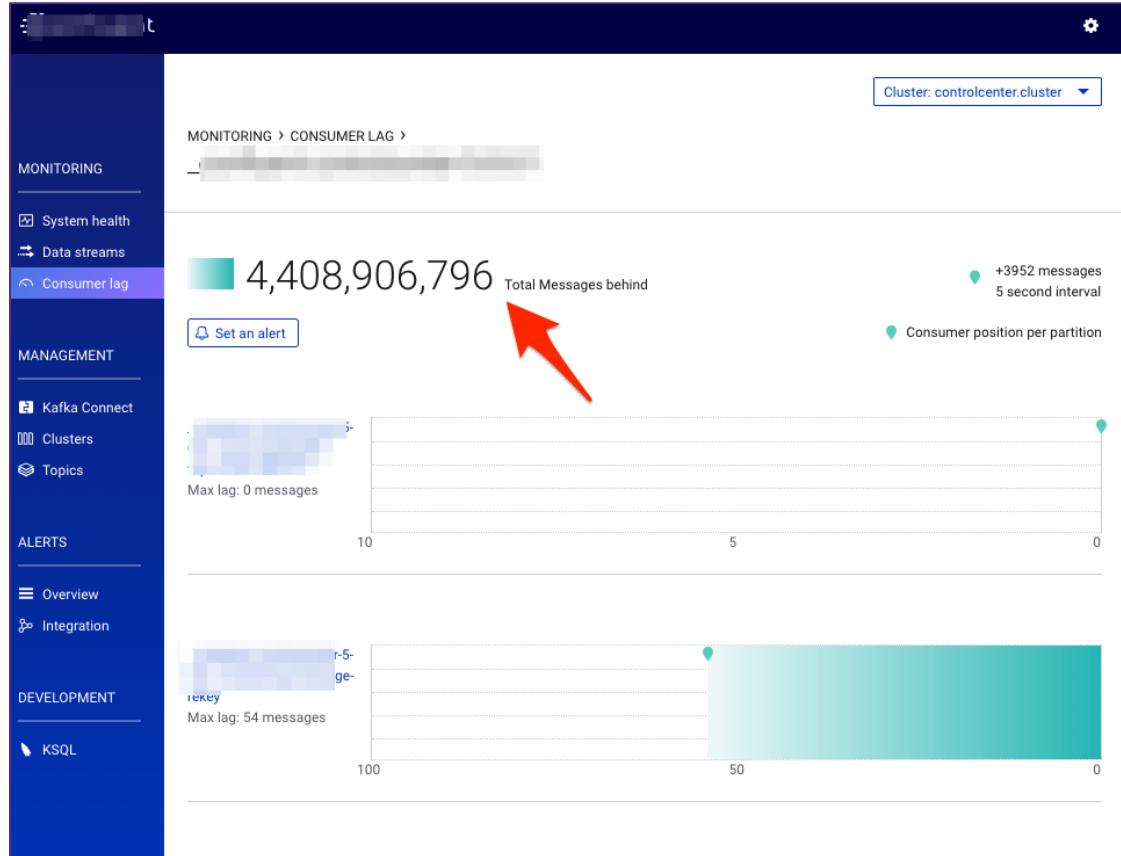
What did we discover?

Slow to fast ... to *really* fast!



Challenge

Lots of data



Hot code #1

A bit of caching

Before: very slow transform
30 records / sec / table

```
-      doApply = transform[R](new AvroData(new AvroDataConfig(props)), schemaFetcher, registerSchema)
+      doApply = transform[R](new AvroData(new AvroDataConfig(props)), cachingSchemaFetcher, registerSchema)
```

After: add cache for schema lookup
200 records / sec / table

```
val cachingSchemaFetcher: ConnectRecord[R] => Throwables \/ Schema =
  record => schemaCache.get(record.topic()) match {
    case Some(schema) => schema.right[Throwables]
    case None =>
      val result = for {
        initTopicName <- \/.fromEither(TopicName.fetch(record.topic))
        .leftMap(_ => new IllegalArgumentException(s"Invalid init topic format: ${record.topic}"))
        schema <- schemaFetcher.apply(initTopicName)
      } yield schema
      result.foreach(schema => schemaCache = schemaCache + (record.topic() -> schema))
      result
  }
```

Hot code #2

A bit more caching

Before: still slow transform
200 records / sec / table

```
doApply = transform[R](new AvroData(new AvroDataConfig(props)), cachingSchemaFetcher, registerSchema)
doApply = transform[R](new AvroData(new AvroDataConfig(props)), cachingSchemaFetcher, registerSchema, cachingDocParser)
```

After: add cache for field metadata
5,500 records / sec / table

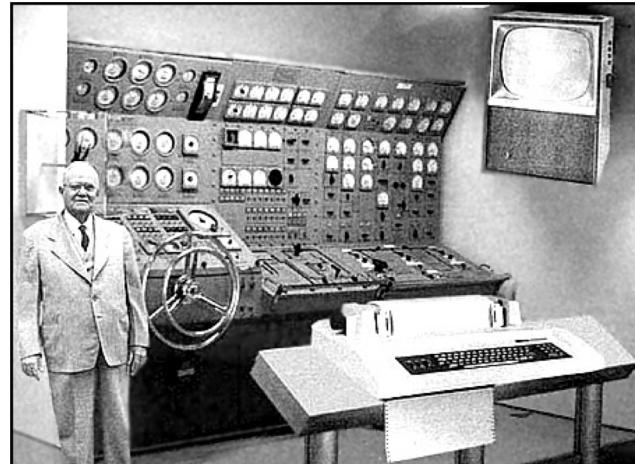
```
val cachingDocParser: String => String \> PwxAvroFieldDocComment =
  docField => docCommentCache.get(docField) match {
    case Some(docComment) => docComment.right[String]
    case None =>
      val result = PwxAvroFieldDocComment.parse(docField)
      result.foreach(docComment => docCommentCache = docCommentCache + (docField -> docComment))
      result
  }
```

Horizontal scaling?

Theory

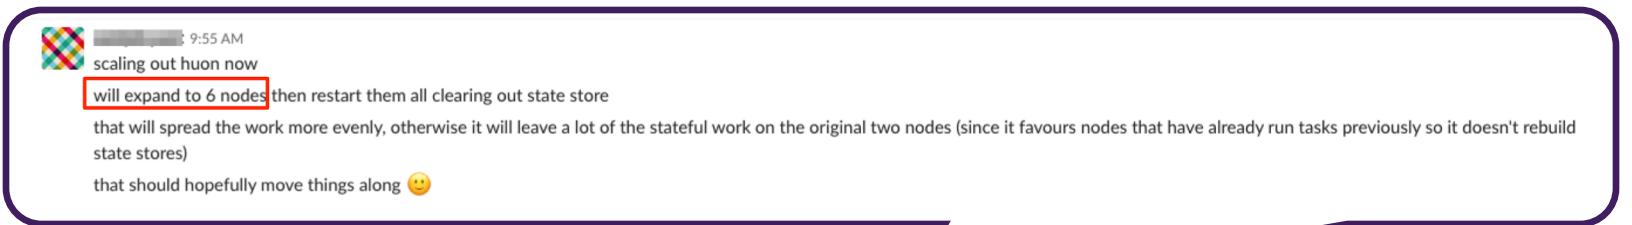
To scale out, you simply start another instance of your stream processing application, e.g. on another machine. The instances of your application will become aware of each other and automatically begin to share the processing work.

<https://www.confluent.io/blog/elastic-scaling-in-kafka-streams/>



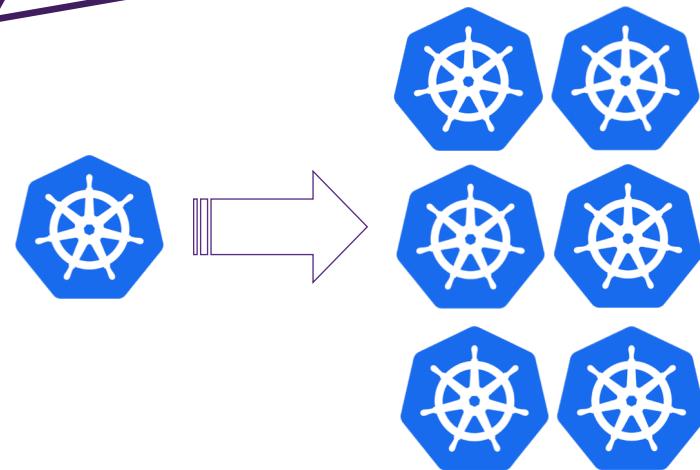
Horizontal scaling ... scales horizontally!

Testing



A screenshot of the DataEng service configuration interface. The service selected is 'huon-motorvehicle' in the stack 'cld-sva-transform'. The configuration fields shown are:

- Type: Service
- Scale: 1 (with a red arrow pointing to the plus sign button)
- Image: swrepos.ai [redacted] d-sva-vehicle-transformer:latest
- Entrypoint: [redacted]



Horizontal scaling ... meet efficient code

Reality

30 records / sec
Approx 100,000 / hr



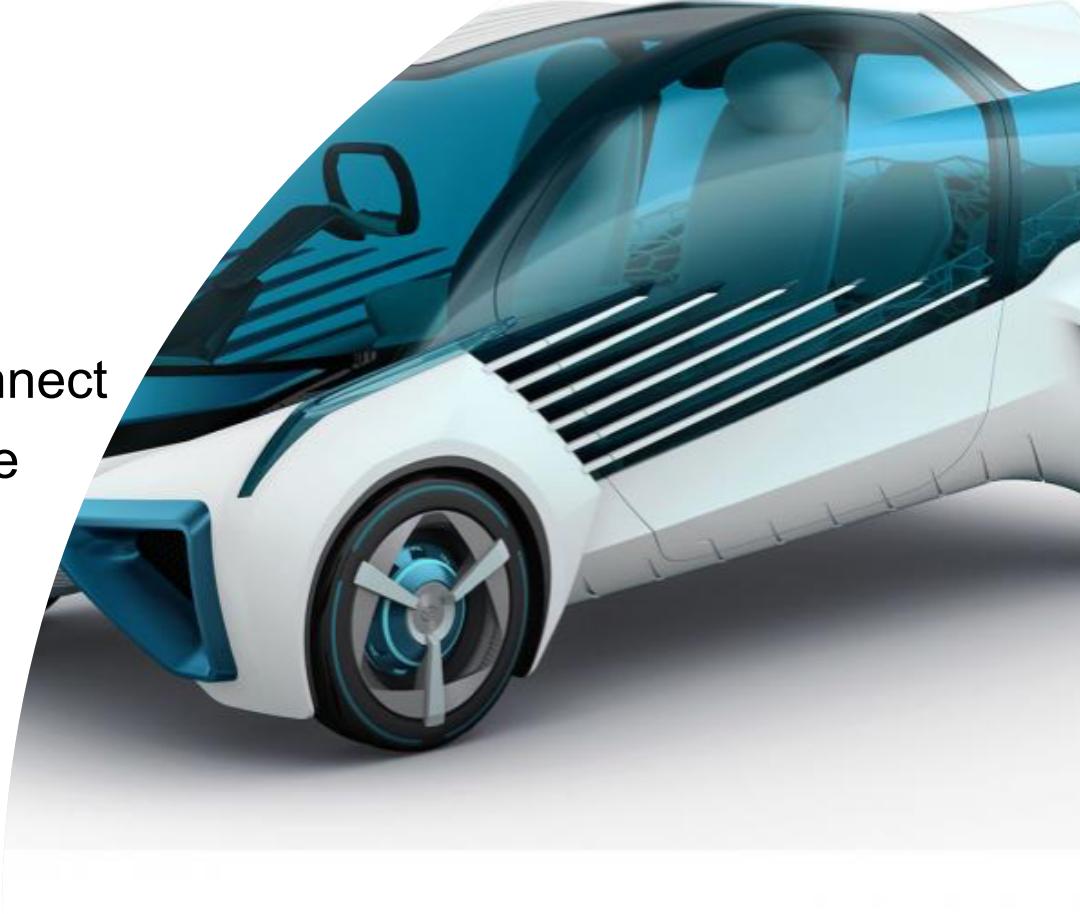
33,500 records / sec
Approx 130 mill / hr



Summary

What did we cover again?

- Architecture of our data flow
 - Extract - Low touch CDC
 - Transformation & Match
 - Serving ; Kafka & Kafka Connect
- Solutions for high performance





Questions?



New Zealand



Australia



Asia